

CONTENTS:

- 1) Distinct Vs Group By Functions
- 2) Explain
- 3) Format Functions
- 4) Help and Show
- 5) Join Functions
- 6) Join Indexes

Distinct Vs. Group By

Distinct Vs. Group By**The Distinct Command**

| Student Table | | | | |
|---------------|-----------|------------|------------|----------|
| Student_ID | Last_Name | First_Name | Class_Code | Grade_Pt |
| 423400 | Larkins | Michael | FR | 0.00 |
| 125624 | Hanson | Henry | FR | 2.88 |
| 200020 | McRoberts | Richard | JR | 1.90 |
| 240000 | Johnson | Stanley | ? | ? |
| 231222 | Wilson | Suzie | SO | 3.80 |
| 234121 | Thomas | Wendy | FR | 4.00 |
| 324452 | DeLaney | Sanny | SR | 3.25 |
| 120350 | Phillips | Martin | SR | 3.00 |
| 322130 | Bond | Jimmy | JR | 3.95 |
| 332450 | Smith | Andy | SO | 2.00 |

| Class_Code |
|------------|
| ? |
| FR |
| SO |
| JR |
| SR |

```
SELECT Distinct Class_Code
FROM Student_Table
ORDER BY 1;
```

DISTINCT eliminates duplicates from returning in the Answer Set.

Distinct vs. GROUP BY

| Student Table | | | | |
|---------------|-----------|------------|------------|----------|
| Student_ID | Last_Name | First_Name | Class_Code | Grade_Pt |
| 423400 | Larkins | Michael | FR | 0.00 |
| 125624 | Hanson | Henry | FR | 2.88 |
| 200020 | McRoberts | Richard | JR | 1.90 |
| 240000 | Johnson | Stanley | ? | ? |
| 231222 | Wilson | Suzie | SO | 3.80 |
| 234121 | Thomas | Wendy | FR | 4.00 |
| 324452 | DeLaney | Sanny | SR | 3.25 |
| 120350 | Phillips | Martin | SR | 3.00 |
| 322130 | Bond | Jimmy | JR | 3.95 |
| 332450 | Smith | Andy | SO | 2.00 |

```
SELECT Distinct Class_Code
FROM Student_Table
ORDER BY 1;
```

```
SELECT Class_Code
FROM Student_Table
GROUP BY 1;
```

| Class_Code |
|------------|
| ? |
| FR |
| SO |
| JR |
| SR |

DISTINCT and GROUP BY in the two examples return the same answer set.

Rules of Thumb for DISTINCT vs. GROUP BY

```
SELECT DISTINCT Class_Code
FROM Student_Table
ORDER BY 1;
```

```
SELECT Class_Code
FROM Student_Table
GROUP BY 1
ORDER BY 1;
```

Rules for DISTINCT Vs. GROUP BY

1. Many Duplicates – use GROUP BY
2. Few Duplicates – use DISTINCT
3. Space Extended – use GROUP BY

DISTINCT and GROUP BY work similarly, but utilize both for different situations.

GROUP BY Vs. DISTINCT – Good Advice

It is better to use this example

```
Select First_Name, Last_Name From Employee_Table
Group By 1,2;
```

This example could cause skewing

```
Select DISTINCT First_Name, Last_Name From
Employee_Table;
```

Ad Hoc users should use the GROUP BY example above instead of the DISTINCT example. Both queries will provide the same results, but the GROUP BY could be significantly faster. Teradata has improved this issue with Teradata V13, but in some cases, the DISTINCT still causes skewing and is slower.

Quiz – How many rows come back from the Distinct?

| Student_ID | Student_Table | | Class_Code | Grade_Pt |
|------------|---------------|------------|------------|----------|
| | Last_Name | First_Name | | |
| 023400 | Lackins | Michael | PS | 0.00 |
| 023404 | Amerson | Henry | PS | 2.00 |
| 020023 | Hobbsworth | Richard | PS | 1.00 |
| 020000 | Johnson | Stanley | ? | 0 |
| 021222 | Wilson | Stacie | SO | 3.80 |
| 024121 | Thomas | Wendy | PS | 4.00 |
| 024401 | Delaney | Deany | SR | 3.35 |
| 021050 | Phillips | Wanda | SR | 3.00 |
| 021110 | Wood | Jimmy | SR | 3.65 |
| 021400 | Smith | Andy | SO | 2.00 |

```
SELECT
  Distinct Class_Code, Grade_Pt
FROM
  Student_Table
ORDER BY
  Class_Code, Grade_Pt;
```

How many rows will come back from the above SQL?

Answer - How many rows come back from the Distinct?

```
SELECT DISTINCT Class_Code, Grade_Pt  
FROM Student_Table  
ORDER BY Class_Code, Grade_Pt
```

| Class_Code | Grade_Pt |
|------------|----------|
| F | 0 |
| F2 | 0.00 |
| F2 | 2.00 |
| F2 | 4.00 |
| F2 | 1.00 |
| F2 | 2.50 |
| F2 | 2.00 |
| F2 | 2.00 |
| F2 | 3.00 |
| F2 | 3.00 |

No Rows have
the exact same
Class_Code
and
Grade_Pt Each
row is distinct

How many rows will come back from the above SQL? 10. All rows came back. Why? Because there are no exact duplicates that contain a duplicate Class_Code and Duplicate Grade_Pt combined. Each row in the SELECT list is distinct.

Explain

Explain

EXPLAIN Keywords

| | |
|----------------------------|---|
| Locking Pseudo Table | Serial lock on a symbolic table. Every table has one. Used to prevent deadlocks situations between users. |
| Locking table for | Indicates that an ACCESS, READ, WRITE, or EXCLUSIVE lock has been placed on the table. |
| Locking rows for <type> | Indicates that an ACCESS, READ, or WRITE lock is placed on rows read or written. |
| On an ABOUT read | Guarantees a transaction is not in progress for the user. |
| All AMPs retrieve | All AMPs are receiving the AMP stops and are involved in providing the answer set. |
| By way of an all rows scan | Rows are read sequentially on all AMPs. |
| By way of primary index | Rows are read using the Primary Index (columns). |
| By way of index number | Rows are read using the Secondary index – number from HELP INDEX. |
| BMMS | Bit Map Set Manipulation Step, alternative direct access technique when multiple NUS columns are referenced in the WHERE clause. |
| Residual conditions | WHERE clause conditions, other than those of a join. |
| Eliminating duplicate rows | Providing unique values, normally result of DISTINCT, GROUP BY or subquery. |
| Where unknown | Indicates that NULL values will not compare to a TRUE or FALSE. Seen in a subquery using comparison will be ignored. |
| Not in or NOT - ALL | Indicates no rows will be returned or ignored comparison. |
| Nested join | The fastest join possible. It uses a LPI to retrieve a single row after using a LPI or a USI as the WHERE to reduce the join to a single row. |

EXPLAIN Keywords Continued

| | |
|--|---|
| Merge join | Rows of one table are matched to the other table on common domain columns after having sorted into the same sequence, normally from Hash. |
| Product join | Rows of one table are matched to all rows of another table with no concern for domain match. |
| ROWID join | A very fast join. It uses the ROWID of a LPI to retrieve a single row after using a LPI or a USI in the WHERE to reduce the join to a single row. |
| Duplicated on all AMPs | Participating rows for the table (normally smaller tables) of a join are duplicated on all AMPs. |
| Hash redistributed on all AMPs | Participating rows of a join are hashed on the join column and sent to the same AMP that stores the matching row of the table to join. |
| SDMS | Set Manipulation Step, result of an INTERSECT, UNION, EXCEPT or MINUS operation. |
| Last use | SPOOL file is no longer needed after the step and space is released. |
| Built locally on the AMPs | As rows are read, they are put into SPOOL on the same AMP. |
| Aggregate Intermediate Results computed locally | The aggregation values are all on the same AMP and therefore no need to redistribute them to work with rows on other AMPs. |
| Aggregate Intermediate Results computed globally | The aggregation values are not all on the same AMP and must be redistributed on one AMP to accompany the same value with from the other AMPs. |

Explain Example – Full Table Scan

```
EXPLAIN SELECT * FROM Employee_Table;
```

1. First, we lock a distinct SQL_CLASS "pseudo table" for read on a RowId hash to prevent global deadlock for SQL_CLASS Employee_Table.
2. Next, we lock SQL_CLASS Employee_Table for read.
3. We do an all-AMPs RETRIEVE step from SQL_CLASS Employee_Table by way of an all-rows scan with no residual conditions into Spool 1 (group_amps), which is built locally on the AMPs. The size of Spool 1 is estimated with low confidence to be 6 rows (642 bytes). The estimated time for this step is 0.03 seconds.
4. Finally, we send out an END TRANSACTION step to all AMPs involved in processing the request.
 - > The contents of Spool 1 are sent back to the user as the result of statement 1. The total estimated time is 0.03 seconds.

When you use all-AMPs RETRIEVE by way of an all-rows scan, that means that Teradata is doing a Full Table Scan, and thus it is reading every row in the table. All AMPs are retrieving all the rows they own.

Explain Example - Unique Primary Index (UPI)

```
EXPLAIN SELECT * FROM Employee_Table
WHERE Employee_No = 2000000;
```

1. First, we do a single-AMP RETRIEVE step from SQL_CLASS Employee_Table by way of the unique primary index "SQL_CLASS Employee_Table Employee_No = 2000000" with no residual conditions. The estimated time for this step is 0.01 seconds.
 - > The row is sent directly back to the user as the result of statement 1. The total estimated time is 0.01 seconds.

If you use the Primary Index column in the WHERE clause, you will most likely get a Single-AMP retrieve by way of the Unique Primary Index. This is the fastest query!

Explain Example - Non-Unique Primary Index (NUPI)

```
EXPLAIN SELECT * FROM Sales_Table
WHERE Product_ID = 1000;
```

1. First, we do a single-AMP RETRIEVE step from SQL_CLASS Sales_Table by way of the primary index "SQL_CLASS Sales_Table Product_ID = 1000" with no residual conditions into Spool 1 (one-amp), which is built locally on that AMP. The size of Spool 1 is estimated with low confidence to be 2 rows (160 bytes). The estimated time for this step is 0.02 seconds.
 - > The contents of Spool 1 are sent back to the user as the result of statement 1. The total estimated time is 0.02 seconds.

If you use the Primary Index column in the WHERE clause, you will most likely get a Single-AMP retrieve. This example utilized a Non-Unique Primary Index (NUPI). It doesn't get much faster than this, unless it utilizes a Unique Primary Index (UPI).

Explain Example - Unique Secondary Index (USI)

```
CREATE UNIQUE INDEX (First_Name, Last_Name) on Employee_Table;
EXPLAIN SELECT * FROM Employee_Table
WHERE First_Name = 'Lorraine'
AND Last_Name = 'Larkins';
```

Created a Unique
Secondary Index

- 1) First, we do a two-AMP RETRIEVE step from SQL_CLASS Employee_Table by way of unique index # 12 "SQL_CLASS Employee_Table Last_Name = Larkins".

SQL_CLASS Employee_Table First_Name = 'Loraine' with no residual conditions. The estimated time for this step is 0.01 seconds.

> The row is sent directly back to the user as the result of statement 1. The total estimated time is 0.01 seconds.

Using a UNIQUE Secondary Index (USI) in the WHERE clause will result in a two-AMP Retrieve every time. This is very fast.

Explain Example – Redistributed to All-AMPs

```
EXPLAIN SELECT E.*, D.*
FROM Employee_Table as E
INNER JOIN
Department_Table as D
ON E.Dept_No = D.Dept_No;
```

- 4) We do an **all-AMPs RETRIEVE** step from SQL_CLASS E by way of an all-rows scan with a condition of ("NOT (SQL_CLASS E.Dept_No IS NULL)") into Spool 2 (all_amps), which is **redistributed** by the **hash** output of (SQL_CLASS E.Dept_No) to all-AMPs. Then we do a SORT to order Spool 2 by row hash. The size of Spool 2 is estimated with low confidence to be 8 rows (204 bytes). The estimated time for this step is 0.01 seconds.
- 5) We do an all-AMPs JOIN step from SQL_CLASS D by way of a Rowhash match scan, which is joined to Spool 2 (last lines) by way of a Rowhash match scan. SQL_CLASS D and Spool 2 are joined using a merge join, with a join condition of ("Dept_No = SQL_CLASS D.Dept_No").

Data is often redistributed on a Join to ensure that the matching rows are on the same physical AMP's in Spool. If you see the word redistributed, then data is being moved!

Explain Example – Row Hash Match Scan

```
EXPLAIN SELECT E1.*, D.*
FROM Employee_Table as E1
INNER JOIN
Department_Table as D
ON E1.Dept_No = D.Dept_No;
```

- 4) We do an all-AMPs JOIN step from SQL_CLASS D by way of a **rowhash match scan**, which is joined to SQL_CLASS E2 by way of a Rowhash match scan. SQL_CLASS D and SQL_CLASS E2 are joined using a merge join, with a join condition of ("SQL_CLASS E2.Dept_No = SQL_CLASS D.Dept_No"). The result goes into Spool 1 (group_amps), which is distributed locally on the AMP's. The size of Spool 1 is estimated with low confidence to be 8 rows (728 bytes). The estimated time for this step is 0.04 seconds.
- 5) Finally, we send out an (END TRANSACTION) step to all AMP's involved in processing the request.

When you use the words Row Hash Match Scan, then the matching rows are on the same physical AMP in Spool and the join takes place. This is what you want to see for joins, and it is even better when data isn't moved beforehand.

Explain Example – Duplicated on All-AMPs

```
EXPLAIN SELECT E.*, D.*
FROM Department_Table as D,
Department_Table as D
WHERE D.Dept_No = D.Dept_No;
```

- 4) We execute the following steps in parallel:
 - 1) We do an all-AMPs RETRIEVE step from SQL_CLASS D by way of an all-rows scan with a condition of ("NOT (SQL_CLASS D.Dept_No IS NULL)") into Spool 2 (all_amps), which is **redistributed** to all-AMPs. Then we do a SORT to order Spool 2 by the hash code of (SQL_CLASS D.Dept_No). The size of Spool 2 is estimated with low confidence to be 16 rows (728 bytes). The estimated time for this step is 0.01 seconds.

Data is often redistributed on a Join to ensure that the matching rows are on the same physical AMP's in Spool. If you see the word redistributed, then data is being moved! This means that the smaller table (usually) is duplicated on each AMP. The Parsing Engine decided that this was a less costly approach than redistributing the data. This is sometimes called a

Big Table/Small Table Join

Explain Example - Low Confidence

```
EXPLAIN SELECT * FROM Addresses;
```

- 1) First, we lock a distinct SQL_CLASS "pseudo table" for read on a Rowlock to prevent global deadlock for SQL_CLASS Addresses.
- 2) Next, we lock SQL_CLASS Addresses for read.
- 3) We do an all-AMPs RETRIEVE step from SQL_CLASS Addresses by way of an all-rows scan with no residual conditions into Spool 1 (group_amps), which is built locally on the AMPs. The size of Spool 1 is estimated with low confidence to be 2 rows (114 bytes). The estimated time for this step is 0.03 seconds.
- 4) Finally, we send out an END TRANSACTION step to all AMPs involved in processing the request.
 - > The contents of Spool 1 are sent back to the user as the result of statement 1. The total estimated time is 0.03 seconds.

The EXPLAIN plan estimates 2 rows, but there are actually 5 rows coming back. When the Explain plan shows low confidence, it is often because there are no COLLECT STATISTICS on the table. So, the PE estimates. Statistics collection is often done by the DBA. The next slide will COLLECT STATISTICS and retry the query.

Explain Example - High Confidence

```
1 COLLECT STATISTICS ON Addresses
   COLUMN Subscribers_Pk;
2 EXPLAIN SELECT * FROM Addresses;
```

- 3) We do an all-AMPs RETRIEVE step from SQL_CLASS Addresses by way of an all-rows scan with no residual conditions into Spool 1 (group_amps), which is built locally on the AMPs. The size of Spool 1 is estimated with high confidence to be 5 rows (260 bytes). The estimated time for this step is 0.03 seconds.
- 4) Finally, we send out an END TRANSACTION step to all AMPs involved in processing the request.

The EXPLAIN plan now estimates 5 rows and shows High confidence because there are COLLECT STATISTICS on the table. If No Statistics are on the table, the Parsing Engine will make an estimate after sampling a Random AMP.

Explain Example - Product Join

```
EXPLAIN SELECT *
FROM Student_Table S, Course_Table C, Student_Course_Table SC
WHERE S.StudID = SC.SID AND C.CourseID = SC.CID;
```

- 6) We do an all-AMPs JOIN step from SQL_CLASS C by way of an all-rows scan with no residual conditions, which is joined to Spool 2 (last) by way of an all-rows scan. SQL_CLASS C and Spool 2 are joined using a product join, with join condition of (1)=1. The result goes into Spool 1 (group_amps), which is built locally on the AMPs. The size of Spool 1 is estimated with low confidence to be 96 rows (7,564 bytes). The estimated time for this step is 0.05 seconds.

The product join in step 6 is using (1=1) as the join condition where it should be a merge join. Therefore, this is a Cartesian product join. A careful analysis of the SELECT shows a single join condition in the WHERE clause. However, this is a three-table join and should have two join conditions. The WHERE clause needs to be fixed, and by using the EXPLAIN, we have saved valuable time.

Explain Example - BMSMS

```
EXPLAIN SELECT * FROM Employee_Table as E
WHERE Salary > 10000.00
AND Dept_ID = 400;
```

3) We do an **BMSMS (Bit Map Set Manipulation)** step that builds a bit map for Employee_Table by way of index #4 Salary = 30000.00 which is placed in Spool 2. The estimated time for this step is 0.01 seconds.

4) We do an **all-AMPS RETRIEVE** step from 1 by way of index #4 Salary. New AMP and the bit map in Spool 2 is used with a residual condition of (E.Salary = 30000.00) into Spool 1 (group_amps) which is built locally on the AMPs. The size of Spool 1 is estimated with low confidence to be 50 rows (400 bytes). The estimated time for this step is 0.03 seconds.

BMSMS (Bit Map Set Manipulation Step) is an excellent way to process large tables. This basically occurs when multiple columns are ANDed together with each Column being a Non-Unique Secondary Index (NUSI). This usually won't happen unless STATISTICS were collected on the table.

Explain Terminology for Partitioned Primary Index Tables

"**A single partition of**" means that an AMP will access a single partition of a table. This is called Partition Elimination. Only a small slice of the table is read.

"**N partitions of**" means that an AMP will access N partitions of a table. This is called also considered Partition Elimination. The smaller the number N is, the faster the query will be compared to the usual Full Table Scan.

"**SORT to partition Spool n by Rowkey**" means the spool is to be sorted by Rowkey, which constitutes the Partition Number and the Hash of the Primary Index. This is done usually for a Join.

"**Rowkey based**" means an equality join on the Rowkey, which is similar to the Row Hash Match Scan for Non-PPH Tables. The join is done prior.

"**Eliminated by dynamic partition**" means a join condition where dynamic partition elimination has been used.

Partitioned Primary Index (PPH) tables are tables that are not sorted by Row Hash on each AMP, but instead sorted first by the Partition. This is designed to eliminate full table scans on Range Queries. When the explain shows Partition Elimination, then a Full Table Scan is NOT being performed, which is the entire purpose of PPH Tables.

Explain Example - From a Single Partition

```
EXPLAIN SELECT * FROM Order_PPH
WHERE Order_Date = '1998-05-04' ;
```

3) We do an **all-AMPS RETRIEVE** step from a **single partition** of SQL_CLASS_Order_PPH with a condition of (SQL_CLASS_Order_PPH.Order_Date = DATE '1998-05-04') and a residual condition of (SQL_CLASS_Order_PPH.Order_Date = DATE '1998-05-04') into Spool 1 (group_amps) which is built locally on the AMPs. The size of Spool 1 is estimated with no confidence to be 1 row (41 bytes). The estimated time for this step is 0.02 seconds.

The above example uses all-AMPs, but each AMP only reads a single partition. This is as good as it gets with a Partition Primary Index (PPH) Table.

Explain Example - From N Partitions

```
EXPLAIN SELECT * FROM Order_PPH
WHERE Order_Date BETWEEN '1998-05-01' AND '1998-05-31' ;
```

3) We do an **all-AMPS RETRIEVE** step from **31 partitions** of SQL_CLASS_Order_PPH with a condition of (SQL_CLASS_Order_PPH.Order_Date = DATE '1998-05-31') AND (SQL_CLASS_Order_PPH.Order_Date = DATE '1998-05-01') into Spool 1 (group_amps) which is built locally on the AMPs. The size of Spool 1 is estimated with no confidence to be 2 rows (82 bytes). The estimated time for this step is 0.03 seconds.

4) Finally, we send out an END TRANSACTION step to all AMPs involved in processing the request.

The above example uses all-AMPs, but each AMP only 31 partitions. This has eliminated reading the entire table, and dramatically improved the performance on Range Queries like the example above.

Explain Example - Partitions and Current Date

```
RECALCIN SELECT * FROM Order_FPI  
WHERE Order_Date = Current_Date )
```

- 3) We do an **all-AMPs** **RECALCIN** step from a single condition of SQL_CLASS Order 270 with a condition of (SQL_CLASS Order_FPI Order_Date = DATE '2012-06-18') with a residual condition of (SQL_CLASS Order_FPI Order_Date = DATE '2012-06-18') into Spool 1 (global arena), which is built locally on the AMPs. The size of Spool 1 is estimated with no confidence to be 1 row (41 bytes). The estimated time for this step is 0.00 seconds.

The above example uses all-AMPs, but each AMP reads only one partition. What is different about this is the improvement that Teradata has made to Current_Date. This happened in Teradata V12 and is a welcomed addition.

Format Functions

Format Functions

The FORMAT Command

```
SELECT Current_Date (FORMAT 'mm-dd-yy') ;
```

DATE
05-07-13

Today's date has
been formatted
for a 2-digit 'year'

In this example, we are using it for dates. All dates in Teradata are stored in the systems as an INTERGERDATE. This allows it to be read and formatted easily.

The Basics of the FORMAT Command

```
SELECT Current_Date (FORMAT 'mm=dd-yy');
```

When you put 'mm', you will get the month as a two digit number.

```
SELECT Current_Date (FORMAT 'mm=dd-yy');
```

When you put 'dd', you are formatting the day as a two digit number.

```
SELECT Current_Date (FORMAT 'mm=dd-yy');
```

When you put 'yy', you are formatting the year as a two digit number.

Format the dates for appearance on the report. The actual data doesn't change.

Quiz - How will the Date Appear after Formatting

```
Current_Date = Month_2013_2013  
SELECT Current_Date (FORMAT 'mm=dd-yy') ;  
How will the answer appear?
```



Answer to Quiz – How will the Date Appear after Formatting

```
Current_Date = March 20th, 2018  
SELECT Current_Date (FORMAT'mm-dd-yyyy') ;  
How will the answer appear?  
ANSWER: 03-20-2018
```

Quiz – How will the Date Appear after Formatting

```
Current_Date = March 20th, 2018  
SELECT Current_Date (FORMAT'mm-dd-yyyy') ;  
How will the answer appear?
```



Answer to Quiz – How will the Date Appear after Formatting

```
Current_Date = March 20th, 2018  
SELECT Current_Date (FORMAT'mm-dd-yyyy') ;  
How will the answer appear?  
ANSWER: 03-20-2018
```

Formatting with MMM for the Abbreviated Month

```
Current_Date = March 20th, 2018  
SELECT Current_Date (FORMAT'mmm-dd-yyyy') ;  
How will the answer appear?
```



Answer to Quiz – How will the Date Appear after Formatting

```
Current_Date = March 20th, 2018
```



```
SELECT Current_Date (FORMAT 'mmmm-dd-yyyy');
```



How will the answer appear?
You get the first three letters of the month!

```
ANSWER: Mar-20-2013
```

Formatting with MMMM for the Full Month Name

```
Current_Date = March_20th, 2013
```

```
SELECT Current_Date (FORMAT 'mmmm-dd-yyyy');
```



How will the answer appear?



Formatting with MMMM for the Full Month

```
Current_Date = March_20th, 2013
```

```
SELECT Current_Date (FORMAT 'mmmm-dd-yyyy');
```

How will the answer appear?

ANSWER: March-20-2013

Formatting with DDD for the Julian Day

```
Current_Date = March 20th, 2013
```

```
SELECT Current_Date (FORMAT 'mm-ddd-yyyy');
```

How will the answer appear?

ANSWER: 03-079-2013

Julian Date

Formatting with DDD for the Julian Day

```
Current_Date = March 20th, 2013
```

```
SELECT Current_Date (FORMAT 'mm-ddd-yyyy');
```

How will the answer appear?

?

Formatting with EEE or EEEE for the Day of the Week

```
Current_Date = March 20th, 2013
```

```
SELECT Current_Date (FORMAT 'eee-mm-dd-yyyy');
```

There are 3 e's in front!

```
SELECT Current_Date (FORMAT 'ccc-mm-dd-yyyy');
```

There are 4 e's in front!

How will the answers appear?

?

EEEE for the Abbreviated or Full Day of the Week

Current Date = March 20th, 2013

How will the answers appear?

```
SELECT Current_Date (FORMAT 'ccc-mm-dd-yyyy');
```

ANSWER: Sun-03-20-13

There are 3 e's in front!

```
SELECT Current_Date (FORMAT 'ccc-mm-dd-yyyy');
```

ANSWER: Sunday-03-20-13

There are 4 e's in front!

Placing Spaces inside your Formatting Commands with a B

Current Date = March 20th, 2013

```
SELECT Current_Date (FORMAT 'ccccB000000');
```

How will the answers appear?

?

Formatting Spaces with B or b

Current Date = March 20th, 2011

SELECT Current_Date (FORMAT 'eeeeBBbbbMMDD');

↑↑↑↑
(4 Blank Spaces)

How will the answer appear?

?

ANSWER: Sunday March

Formatting with 9

You can also format how numbers appear. Such as in the case of a phone number.

```
SELECT 5133000346 (FORMAT '999-999-9999');
```

How will the answer appear?

?

By putting in 999-999-9999, this is telling the system to put the literal numbers 5133000341 next to the SELECT into the formatting style.

Formatting with 9 Results

You can also format how numbers appear. Such as in the case of a phone number.

```
SELECT 5133000346 (FORMAT '999-999-9999');
```

How will the answer appear?

513-300-0346

By putting in 999-999-9999, this is telling the system to put the literal numbers 5133000341 next to the SELECT into the formatting style.

Troubleshooting when Formatted Data Overflows

Notice that we've taken out one of the 9s in the Format Statement.

```
SELECT 5133000349 (FORMAT '99-999-9999');
```

How will the answer appear?

?

The **FORMAT OVERFLOW** is what happens when your system doesn't have enough spaces in the **FORMAT** to cover the number you are trying to format.

Troubleshooting when Formatted Data Overflows

Notice that we've taken out one of the **9s** in the Format Statement

```
SELECT '5120000340' (FORMAT '99-999-9999');
```

This is not an error, but
something is "wrong".

The **FORMAT OVERFLOW** is what happens when your system doesn't have enough spaces in the **FORMAT** to cover the number you are trying to format.

Formatting with X or x

You can also format letters and words!

```
SELECT 'ABCDE' (FORMAT 'X5');
```

How will the answer appear?

?

You can also **FORMAT** characters. Look at this example. It doesn't matter if the **X's** are capitalized or not.

Formatting with Z

The **Z's** represent potential data. This tells the system that if there is a number to put in the **Z's** position, then put it in. If there is not, leave it blank.

```
SELECT 1021.53 (FORMAT 'ZZZZZ9.99');
```

How will the answer appear?

?

A 'V' in a format statement means that if there is a number in the 'V' position, then put it in. If there isn't one, then you put a blank.

Formatting with Z Visual

The Zs represent potential data. This tells the system that if there is a number to put in the Z's position, then put it in. If there is not, leave it blank.

```
SELECT 1021.53 (FORMAT 'ZZZZZ9.99');
```

How will the answer appear?

1021.53

What a 'V' represents in a format statement is that if there is a number in the 'V' position, then put it in. If there isn't one, then you put a blank.

Formatting with 9

The 9s represent potential data. This tells the system that if there is a number to put in the 9's position, then put it in. If there is not, leave it blank.

```
SELECT 1021.53 (FORMAT '99999999.9999');
```

How will the answer appear?

?

Formatting with \$ Visual

The \$s represent potential data. This tells the system that if there is a number to put in the \$s position, then put it in. If there is not, leave it blank.

```
SELECT 1021.53 (FORMAT '9999999.9999');
```

How will the answer appear?

ANSWER: 00001021.5300

Formatting with \$

What the \$ allows you to do is to tell your formatting to place a \$ sign in front of the result set, but only at the beginning.

```
SELECT 1021.53 (FORMAT 'SSSSSS9.99');
```

How will the answer appear?

?

Formatting with \$ Visual

What the \$ allows you to do is to tell your formatting to place a \$ sign in front of the result set, but only at the beginning.


```
SELECT 1021.53 (FORMAT '$$$$$9.99');
```



How will the answer appear?

ANSWER: \$1021.53

Formatting with \$ and Commas

You can also use commas in your Formatting statements.

```
SELECT 1021.53 (FORMAT '$,$$$,$9.99');
```



How will the answer appear?

**Formatting with \$ and Commas Visual**

You can also use commas in your Formatting statements.

```
SELECT 1021.53 (FORMAT '$,$$$,$$9.99');
```



How will the answer appear?

ANSWER: \$1,021.53

Formatting with \$ and Commas and 9

You can also use commas in your Formatting statements.

```
SELECT 0.53 (FORMAT '$,$$$,$$9.99');
```



How will the answer appear?



Formatting with \$ and Commas and 0 with Zero Digits

You can also use commas in your Formatting statements.

```
SELECT 0.53 (FORMAT '$,$$$,$$9.99');
```



How will the answer appear?

```
ANSWER: $0.53
```

The 'Y' sees the 0 and knows to bring it back in the answer set. The floating \$ will only bring back a \$ for the first character.

A Great Formatting Example

```
SELECT 'NAME' AS Pst_Name,
       '2014050909' AS Pst_Phone,
       '12345' AS Pst_State,
       '991001' AS Pst_Zip,
       '991001' AS Pst_City,
       '991001' AS Pst_Street,
       '991001' AS Pst_Region;
```

```
NAME Pst_Phone Pst_State Pst_City Pst_Street Pst_Region
-----
```

There are only two things that need to be watched when using the FORMAT function. First, the data type must match the formatting character used or a syntax error is returned. So, if the data is numeric, use a numeric formatting character and the same condition for character data. The other concern is configuring the format mask big enough for the largest data column. If the mask is too short, the SQL command executes, however, the output contains a series of "*****" to indicate a format overflow.

A Great Formatting Example for Day, Month, and Year

```
SELECT Current_Date (FORMAT 'mm-dd-yyyy') AS Digit1_18,
       Current_Date (FORMAT 'mm-dd-yyyy') AS Digit1_18;
```



```
SELECT 'They match' as "Do They?"  
WHERE A(CASESPECIFIC) = A(CS);
```

Do They

They did NOT match using the
CASESPECIFIC Command

```
SELECT 'They match' as "Do They?"  
WHERE A = A;
```

Do They

They Match

They match because Teradata Mode
does not differentiate about case

In ANSI Mode, a capital 'A' is seen as different in comparison to a little 'a'. The NOT CASESPECIFIC command, which is abbreviated as NOT CS will make sure the case of a letter is examined and returned whether or not they have the same case.

Using the LOWER Command

```
SELECT LOWER('AbCdE') as Result1;
```

Result1

abcde

All Capital letters are
now in LOWER Case

```
SELECT 'They match' as "Do They?"  
WHERE LOWER(ABCODE) = 'abcde';
```

Do They

They Match

They match because we used
the LOWER command so now
they match perfectly in ANSI Mode

In ANSI Mode, a capital 'A' is seen as different in comparison to a little 'a'. The LOWER command can be used to make sure the case of a column is lowered.

Using the UPPER Command

```
SELECT UPPER('ABCDE') as Result1;
```

Result1
ABCDE

All letters are now
in UPPER Case.

```
SELECT 'They match!' as "Do They"  
WHERE 'ABCDE' = UPPER('abcde');
```

Do They
They Match

They match because we used
the UPPER command, so now
they match perfectly in ANSI Mode.

In ANSI Mode, a capital 'K' is seen as different in comparison as a little 'k'. The UPPER command can be used to make sure the case of a column is UPPERED.

HELP and SHOW

HELP and SHOW

Determining the Release of your Teradata System

```

SELECT * FROM dbc dbcinfo;

dbcinfo
-----
release      13.00.00.12
version      13.00.00.12
language     Standard

```

The above query pulls information from the Data Dictionary in USER DBC. Some companies don't allow users to see this information, but if you have the access rights, then you can run the above query.

Basic HELP Commands

| | |
|--|--|
| HELP DATABASE <database-name> | Displays the names of all the tables (T), views (V), macros (M), and triggers (G) stored in a database and user-written table comments. |
| HELP USER <user-name> | Displays the names of all the tables (T), views (V), macros (M), and triggers (G) stored in a user area and user-written table comments. |
| HELP TABLE <table-name> | Displays the column names, type identifier, and any user-written comments on the columns within a table. |
| HELP VOLATILE TABLE | Displays the names of all Volatile temporary tables active for the user session. |
| HELP VIEW <view-name> | Displays the column names, type identifier, and any user-written comments on the columns within a view. |
| HELP MACRO <macro-name> | Displays the characteristics of parameters passed to it at execution time. |
| HELP PROCEDURE <procedure-name> | Displays the characteristics of parameters passed to it at execution time. |
| HELP TRIGGER <trigger-name> | Displays details created for a trigger. No action time and sequence. |
| HELP COLUMN <table-name> * | Displays detail data describing the column level characteristics. |
| HELP COLUMN <table-name> <column-name> * | |

Other HELP Commands

| | |
|--|--|
| HELP INDEX <table-name> | Displays the indexes and their characteristics like unique or non-unique and the column or columns involved in the index. This data is used by the Optimizer to create a plan for SQL. |
| HELP STATISTICS <table-name> | Displays values associated with the data demographics collected on the table. This data is used by the Optimizer to create a plan for SQL. |
| HELP CONSTRAINT <table-name> <constraint-name> | Displays the checks to be made on the data when it is inserted or updated and the columns are involved. |
| HELP SESSION | Displays the user name, account name, login date and time, current database name, collation code set and character set being used, transaction semantics, time zone and character set (SQL). |
| HELP SQL | Displays a list of available SQL commands and functions. |
| HELP SQL <command> | Displays the basic syntax and options for the actual SQL command inserted in place of the <command>. |

| | |
|-----------------------------|--|
| HELP SPL; | Displays a list of available SPL commands. |
| HELP SPL <command> <n> <n>; | Displays the basic syntax and options for the actual SPL command inserted in place of the <command>. |

The above chart shows HELP commands for information on database tables and sessions, as well as SQL and SPL commands.

HELP DATABASE

| HELP_DATABASE_SQL_Class_1 | | | A complete list of KND | |
|---------------------------|------|--|------------------------|--|
| Table/Type/Session name | Kind | Comment | | |
| Aggregate | A | Aggregate function | | |
| AggregateFunction | B | Combined aggregate/analytical function | | |
| ExternalStoredProcedure | E | External stored procedure | | |
| Function | F | Function | | |
| Trigger | G | Trigger | | |
| Method | H | Method | | |
| JoinIndex | I | Join index | | |
| JournalTable | J | Journal table | | |
| Macro | M | Macro | | |
| HashIndex | N | Hash index | | |
| Procedure | P | Procedure | | |
| QueueTable | Q | Queue table | | |
| TableFunction | R | Table function | | |
| OrderedAnalyticalFunction | S | Ordered analytical function | | |
| Table | T | Table | | |
| UDT | U | UDT | | |
| View | V | View | | |
| Authorization | X | Authorization | | |

Not all columns in the HELP Database SQL_Class were shown, just the important ones. The HELP DATABASE command will show you the objects in your database.

HELP USER

| HELP_USER_SQL_1 | | | A complete list of KND | |
|---------------------------|------|--|------------------------|--|
| Table/Type/Session name | Kind | Comment | | |
| Aggregate | A | Aggregate function | | |
| AggregateFunction | B | Combined aggregate and analytical function | | |
| ExternalStoredProcedure | E | External stored procedure | | |
| Function | F | Function | | |
| Trigger | G | Trigger | | |
| Method | H | Method | | |
| JoinIndex | I | Join index | | |
| JournalTable | J | Journal table | | |
| Macro | M | Macro | | |
| HashIndex | N | Hash index | | |
| Procedure | P | Procedure | | |
| QueueTable | Q | Queue table | | |
| TableFunction | R | Table function | | |
| OrderedAnalyticalFunction | S | Ordered analytical function | | |
| Table | T | Table | | |

Not all columns in the `HELP USER` were shown, just the important ones. The `HELP USER` command will show you the objects in a `USER`. `USER` is a keyword!

HELP TABLE

```
HELP Table SQL_Class.Employee_Table ;
```

| Column Name | Type | Comment | Nullable | Format | Title | MaxLength |
|-------------|------|---------|----------|--------|-------|-----------|
| Employee_No | I | ? | Y | -(10)9 | ? | 4 |
| Dept_No | I2 | ? | Y | -(5)9 | ? | 2 |
| Last_Name | CF | ? | Y | X(20) | ? | 20 |
| First_Name | CV | ? | Y | X(12) | ? | 12 |
| Salary | D | ? | Y | —99 | ? | 4 |



I = Integer
I2 = Smallint
CF = Character Fixed
CV = Character Variable
D = Decimal

Not all columns in the HELP TABLE were shown, just the important ones. The HELP TABLE command will show you information about a table.

Adding a Comment to a Table

```
COMMENT ON TABLE SQL_Class.State_Table 'This table holds State';
```

Help Database SQL_Class;

| Table/View/Macro name | Kind | Comment |
|-----------------------|------|------------------------|
| Testemp_Macro | M | ? |
| Subscribers | T | ? |
| Student_Table | T | ? |
| Student_Course_Table | T | ? |
| Stats_Table | T | This table holds Stats |
| Services | T | ? |
| Sales_Table | T | ? |
| Providers | T | ? |
| Order_Table | T | ? |
| Names_View | V | ? |

The above syntax will place a comment on the table

Adding a Comment to a View

CONCISE ON View SQL_VIEWS.Employee_V 'No Salary is shown'

[Help Database SQL_VIEWS](#)

```

Help>Tran>View>Session
+-----+-----+-----+
| Column | Type | Length |
+-----+-----+-----+
| Session | V | 10 |
| Username | V | 30 |
| Password | V | 30 |
| Department | V | 30 |
| Employee | V | 30 |
| Emp_Fob | V | 30 |
| StartDate | V | 30 |
| Job | V | 30 |
| Salary | V | 30 |

```

The above syntax will place a comment on the View.

SELECT SESSION

```
SELECT SESSION;
```

Session

8692

The SELECT Session command will show you the SESSION Number you received when you logged on to Teradata. The Parsing Engine assigned to manage your session tracks you by this session number.

USER Information Functions

```
SELECT Account
,Database
,Session
,USER;
```

```
Account Database Session USER
-----
DB: DB: DB: DB:
```

The Teradata RDBMS (Relational Database Management System) has incorporated into it functions that provide data regarding a user who has performed a login connect to the system. The following functions make that data available to a user for display or storage. Notice the keyword USER.

HELP SESSION

```
Help>Session
```

```

User Account Login Login Current Collation Char Transaction Current Session
Name Name Date Time Database DB ID Database Database Time Zone
DB: DB: 11/26/17 15:25:19 SQL_CAS9 ASCII ASCII Teradata Subsequence 00:00

```

The HELP Session command will show information about your SESSION. Not all columns were shown above, just the most important ones.

HELP SQL

```
Help>SQL
```

```

SQL SQL Command
-----
ABOUT ACCTD, DATE BEGIN LOGGING

```

```
CREATE SET TABLE SQL_CLASS.Employee_Table ,NO FALLBACK,
NO BEFORE JOURNAL,
NO AFTER JOURNAL,
CHECKSUM = DEFAULT
(
Employee_No INTEGER,
Dept_No SMALLINT,
Last_Name CHAR(20) CHARACTER SET LATIN NOT CASESPECIFIC,
First_Name VARCHAR(12) CHARACTER SET LATIN NOT CASESPECIFIC,
Salary DECIMAL(8,2)
```

```
CREATE FRAGMENT INDEX (Employee_Rn)  
INDEX (Last_Name)  
INDEX (Dept_Name);
```

The above syntax will show the Table DDL (Data Definition Language). It is the table CREATE Statement plus any additional changes such as adding an Index.

SHOW View command for View Create Statement

```
SHOW View SQL_VIEWS.Employee_V;
```

```
CREATE VIEW SQL_VIEWS.Employee_V (Emp_No, Name, Phone, Sal, Dept) AS  
SELECT Employee_No,  
Last_Name,  
First_Name,  
Salary,  
Dept_No  
FROM SQL_VIEWS.Employee_Table;
```

The above syntax will show the View's CREATE Statement.

SHOW Macro command for Macro Create Statement

```
SHOW MACRO MY_Macro;
```

```
CREATE MACRO MJ101.MY_Mac (INPARN1 INTEGER, INPARN2 CHAR(10))
AS
  (SELECT DEPT, DAY_OF_WEEK, AVG(SAL)
   FROM SYS.CALENDAR.CALENDAR SC, MYTABLE
   WHERE CALENDAR_DATE = :INPARN2 (DATE, FORMAT 'YYYYMMDD')
   AND DEPT = :INPARN1
   GROUP BY 1,2);
```

The above syntax will show the Macro's CREATE Statement.

SHOW Trigger command for Trigger Create Statement

```
SHOW TRIGGER AVG_SAL_T;
```

```
CREATE TRIGGER MJ1.AVG_SAL_T
AFTER UPDATE OF (SALARY) ON MJ1.EMPLOYEE
REFERENCING OLD AS OLDROW
NEW AS NEWROW
FOR EACH ROW
WHEN (NEWROW.SALARY >
      (SELECT AVG(BUDGET) * .10 (DECIMAL(10,2))
       FROM MJ101.DEPARTMENT))
  (INSERT INTO MJ101.GREATER 10 PERCENT
   (EMP_NUM ,SAL_DATE ,OLDSAL ,NEWSAL ,PERC_OF_BUDGET)
   VALUES (NEWROW.EMP_NUM ,CURRENT_DATE ,OLDROW.SALARY,
           NEWROW.SALARY));
```

The above syntax will show the Trigger's CREATE Statement.

Join Functions

Join Functions

A two-table join using Non-ANSI Syntax

| Customer_Table | | Order_Table | |
|-----------------|---------------------|--------------|-------------|
| Customer_Number | Customer_Name | Order_Number | Order_Total |
| 11111111 | Billy's Best Choice | 121854 | 11111111 |
| 11111111 | Acme Products | 121855 | 11111111 |
| 11111111 | Acme Consulting | 121856 | 11111111 |
| 11111111 | KIT Flamingo | 121857 | 11111111 |
| 11111111 | Databasee B-D | 121858 | 11111111 |

```
SELECT Customer_Table.Customer_Number, Customer_Name, Order_Number, Order_Total
FROM Customer_Table, Order_Table
WHERE Customer_Table.Customer_Number = Order_Table.Customer_Number;
```

A join combines columns on the report from more than one table. The example above joins the Customer_Table and the Order_Table together. The most complicated part of any join is the JOIN CONDITION. The JOIN CONDITION means what Column from each table is a match. In this case, Customer_Number is a match that establishes the relationship, so this join will happen on matching Customer_Number columns.

A two-table join using Non-ANSI Syntax with Table Alias

| Customer_Table | | Order_Table | |
|-----------------|---------------------|--------------|-------------|
| Customer_Number | Customer_Name | Order_Number | Order_Total |
| 11111111 | Billy's Best Choice | 121854 | 11111111 |
| 11111111 | Acme Products | 121855 | 11111111 |
| 11111111 | Acme Consulting | 121856 | 11111111 |
| 11111111 | KIT Flamingo | 121857 | 11111111 |
| 11111111 | Databasee B-D | 121858 | 11111111 |

```
SELECT Cust.Customer_Number, Customer_Name, Order_Number, Order_Total
FROM Customer_Table as Cust, Order_Table as ORD
WHERE Cust.Customer_Number = ORD.Customer_Number;
```

A join combines columns on the report from more than one table. The example above joins the Customer_Table and the Order_Table together. The most complicated part of any join is the JOIN CONDITION. The JOIN CONDITION means what Column from each table is a match. In this case, Customer_Number is a match that establishes the relationship.

Aliases and Fully Qualifying Columns

| Customer_Table | | Order_Table | |
|-----------------|---------------------|--------------|-------------|
| Customer_Number | Customer_Name | Order_Number | Order_Total |
| 11111111 | Billy's Best Choice | 121854 | 11111111 |
| 11111111 | Acme Products | 121855 | 11111111 |
| 11111111 | Acme Consulting | 121856 | 11111111 |
| 11111111 | KIT Flamingo | 121857 | 11111111 |
| 11111111 | Databasee B-D | 121858 | 11111111 |

```

SELECT Cust.Customer_Number,
       Customer_Name, Order_Number,
       Order_Total
FROM   Customer_Table as Cust,
       Order_Table as ORD
WHERE  Cust.Customer_Number
       = ORD.Customer_Number;

```

Annotations:

- Fully Qualified with CUST
- CUST is new Table ALIAS
- ORD is new Table ALIAS

Customer_Number is a column in both the Customer and Order Tables. CUST Customer_Number fully qualifies the column to specifically state we want the Customer_Number from the Customer_Table. That is why we ALIASED the table names, so we could fully qualify any columns in both tables, or else we receive an error!

A two-table join using ANSI Syntax

| Customer_Table | | Order_Table | |
|-----------------|---------------------|--------------|-------------|
| Customer_Number | Customer_Name | Order_Number | Order_Total |
| 11111111 | Billy's Best Choice | 123456 | 11111111 |
| 11111111 | Acme Products | 123456 | 11111111 |
| 11111111 | Acme Consulting | 123456 | 11111111 |
| 11111111 | KIT Plumbing | 123456 | 11111111 |
| 11111111 | Delaware P.C. | 123456 | 11111111 |

```

SELECT Cust.Customer_Number,
       Customer_Name,
       Order_Number,
       Order_Total
FROM   Customer_Table as Cust
INNER JOIN
       Order_Table as ORD
ON      Cust.Customer_Number
       = ORD.Customer_Number;

```

Annotations:

- ON Keyword is used instead of WHERE
- INNER JOIN Keyword replaces the comma

This is the same join as the previous slide except it is using ANSI syntax. Both will return the same rows with the same performance. Rows are joined when the Customer_Number matches on both tables, but non-matches won't return.

Both Queries have the same Results and Performance

| Customer_Table | | Order_Table | |
|-----------------|---------------------|--------------|-------------|
| Customer_Number | Customer_Name | Order_Number | Order_Total |
| 11111111 | Billy's Best Choice | 123456 | 11111111 |
| 11111111 | Acme Products | 123456 | 11111111 |
| 11111111 | Acme Consulting | 123456 | 11111111 |
| 11111111 | KIT Plumbing | 123456 | 11111111 |
| 11111111 | Delaware P.C. | 123456 | 11111111 |

/* Traditional Syntax */

```

SELECT Cust.Customer_Number,
       Customer_Name,
       Order_Number,
       Order_Total
FROM   Customer_Table as Cust,

```

/* ANSI Syntax */

```

SELECT Cust.Customer_Number,
       Customer_Name,
       Order_Number,
       Order_Total
FROM   Customer_Table as Cust

```

```

WHERE Order_Table as ORD
      Ord_Customer_Number
      = Ord_Customer_Number_1;

INNER JOIN
      Order_Table as ORD
ON
      Ord_Customer_Number
      = Ord_Customer_Number_1;

```

Both of these syntax techniques bring back the same result set and have the same performance. The INNER JOIN is considered ANSI. Which one does Outer Joins?

Quiz - Can You Finish the Join Syntax?

| Employee_Table | | | | | Department_Table | |
|----------------|---------|-----------|------------|----------|------------------|------------------|
| Employee_ID | Dept_ID | Last_Name | First_Name | Salary | Dept_ID | Department_Name |
| 1122278 | 100 | Chambers | Headee | 48850.00 | 100 | Marketing |
| 1156249 | 400 | Harrison | Robert | 58500.00 | 200 | Research and Dev |
| 1342118 | 400 | Beilly | William | 58500.00 | 300 | Sales |
| 2312225 | 300 | Sarkins | Lorraine | 40200.00 | 400 | Customer Support |
| 2000000 | 0 | Jones | Empassy | 32500.00 | 500 | Human Resources |
| 1000234 | 10 | Seythe | Richard | 32500.00 | | |
| 1121334 | 400 | Prackling | Cletus | 58500.00 | | |
| 1124887 | 200 | Coffing | Billy | 41880.88 | | |
| 1333454 | 200 | Smith | John | 48000.00 | | |

```

SELECT First_Name, Last_Name,
       Department_Name
FROM   Employee_Table as E
INNER JOIN

```

Finish the join by placing the missing SQL in the proper place!

Answer to Quiz - Can You Finish the Join Syntax?

| Employee_Table | | | | | Department_Table | |
|----------------|---------|-----------|------------|----------|------------------|------------------|
| Employee_ID | Dept_ID | Last_Name | First_Name | Salary | Dept_ID | Department_Name |
| 1122278 | 100 | Chambers | Headee | 48850.00 | 100 | Marketing |
| 1156249 | 400 | Harrison | Robert | 58500.00 | 200 | Research and Dev |
| 1342118 | 400 | Beilly | William | 58500.00 | 300 | Sales |
| 2312225 | 300 | Sarkins | Lorraine | 40200.00 | 400 | Customer Support |
| 2000000 | 0 | Jones | Empassy | 32500.00 | 500 | Human Resources |
| 1000234 | 10 | Seythe | Richard | 32500.00 | | |
| 1121334 | 400 | Prackling | Cletus | 58500.00 | | |
| 1124887 | 200 | Coffing | Billy | 41880.88 | | |
| 1333454 | 200 | Smith | John | 48000.00 | | |

```

       Department_Name
FROM   Employee_Table as E
INNER JOIN
       Department_Table as D
ON     E.Dept_No = D.Dept_No;

```

This query is ready to run.

Quiz - Can You Find the Error?

| Employee_Table | | | | Department_Table | |
|----------------|---------|------------|----------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | Salary | Dept_No | Department_Name |
| 1232570 | 100 | Chambers | 48950.00 | 100 | Marketing |
| 1232496 | 400 | Baillien | 94500.00 | 200 | Research and Dev |
| 2341219 | 400 | Baillien | 94500.00 | 200 | Sales |
| 2341222 | 300 | Larkin | 40200.00 | 400 | Customer Support |
| 2000000 | ? | Jones | 32000.50 | 500 | Human Resources |
| 1000234 | 10 | Myrba | 32000.00 | | |
| 1121334 | 400 | Strickling | 94500.00 | | |
| 1324457 | 200 | Coffing | 41888.88 | | |
| 1323454 | 200 | Smith | 48950.00 | | |

```

SELECT First_Name, Last_Name, Dept_No
FROM Employee_Table as E
INNER JOIN
  Department_Table as D
ON E.Dept_No = D.Dept_No ;

```

This query has an error! Can you find it?

Answer to Quiz – Can You Find the Error?

| Employee_Table | | | | Department_Table | |
|----------------|---------|------------|----------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | Salary | Dept_No | Department_Name |
| 1232570 | 100 | Chambers | 48950.00 | 100 | Marketing |
| 1232496 | 400 | Baillien | 94500.00 | 200 | Research and Dev |
| 2341219 | 400 | Baillien | 94500.00 | 200 | Sales |
| 2341222 | 300 | Larkin | 40200.00 | 400 | Customer Support |
| 2000000 | ? | Jones | 32000.50 | 500 | Human Resources |
| 1000234 | 10 | Myrba | 32000.00 | | |
| 1121334 | 400 | Strickling | 94500.00 | | |
| 1324457 | 200 | Coffing | 41888.88 | | |
| 1323454 | 200 | Smith | 48950.00 | | |

```

SELECT First_Name, Last_Name, E.Dept_No
FROM Employee_Table as E
INNER JOIN
  Department_Table as D
ON E.Dept_No = D.Dept_No ;

```

If a column in the SELECT list is in both tables, you must fully qualify it.

Quiz – Which rows from both tables Won't Return?

| Employee_Table | | | | Department_Table | |
|----------------|---------|------------|----------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | Salary | Dept_No | Department_Name |
| 1232570 | 100 | Chambers | 48950.00 | 100 | Marketing |
| 1232496 | 400 | Baillien | 94500.00 | 200 | Research and Dev |
| 2341219 | 400 | Baillien | 94500.00 | 200 | Sales |
| 2341222 | 300 | Larkin | 40200.00 | 400 | Customer Support |
| 2000000 | ? | Jones | 32000.50 | 500 | Human Resources |
| 1000234 | 10 | Myrba | 32000.00 | | |
| 1121334 | 400 | Strickling | 94500.00 | | |
| 1324457 | 200 | Coffing | 41888.88 | | |
| 1323454 | 200 | Smith | 48950.00 | | |

```

1324487 200 Coffing Billy 41888.88
1333454 200 Smith John 48000.00

SELECT First_Name, Last_Name,
       Department_Name
FROM   Employee_Table as E
INNER JOIN
       Department_Table as D
ON     E.Dept_No = D.Dept_No

```

An Inner Join returns matching rows, but did you know an Outer Join returns both matching rows and non-matching rows?
You will understand soon!

Answer to Quiz – Which rows from both tables Won't Return?

| Employee_Table | | | | | Department_Table | |
|----------------|---------|-----------|------------|----------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
| 1323276 | 100 | Chambers | Madame | 48850.00 | 100 | Marketing |
| 1324349 | 400 | Barrison | Richard | 54500.00 | 200 | Research and Dev |
| 2341218 | 400 | Pelly | William | 36000.00 | 300 | Sales |
| 2312225 | 300 | Lathia | Shakane | 45200.00 | 400 | Customer Support |
| 2000000 | 1 | Jones | Squiggly | 32000.00 | 500 | Human Resources |
| 1000234 | 10 | Seythe | Richard | 32800.00 | | |
| 1121334 | 400 | Strawling | Cletus | 54500.00 | | |
| 1324487 | 200 | Coffing | Billy | 41888.88 | | |
| 1333454 | 200 | Smith | John | 48000.00 | | |

- 1 Squiggly Jones has a NULL Dept_No
- 2 Richard Smythe has an invalid Dept_No 10
- 3 No Employees work in Department 000

The bottom line is that the three rows excluded did not have a matching Dept_No.

LEFT OUTER JOIN

| Employee_Table | | | | | Department_Table | |
|----------------|---------|-----------|------------|----------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
| 1323276 | 100 | Chambers | Madame | 48850.00 | 100 | Marketing |
| 1324349 | 400 | Barrison | Richard | 54500.00 | 200 | Research and Dev |
| 2341218 | 400 | Pelly | William | 36000.00 | 300 | Sales |
| 2312225 | 300 | Lathia | Shakane | 45200.00 | 400 | Customer Support |
| 2000000 | 1 | Jones | Squiggly | 32000.00 | 500 | Human Resources |
| 1000234 | 10 | Seythe | Richard | 32800.00 | | |
| 1121334 | 400 | Strawling | Cletus | 54500.00 | | |
| 1324487 | 200 | Coffing | Billy | 41888.88 | | |
| 1333454 | 200 | Smith | John | 48000.00 | | |

```

SELECT First_Name, Department_Name
FROM   Employee_Table as E
LEFT OUTER JOIN
       Department_Table as D
ON     E.Dept_No = D.Dept_No;

```

1st Table
after FROM
is always the
LEFT Table

This is a LEFT OUTER JOIN. That means that all rows from the LEFT Table will appear in the report regardless if it finds a match on the right table.

LEFT OUTER JOIN Brings Back All Rows in the Left Table

| Employee_Table | | | | | Department_Table | |
|----------------|---------|-----------|------------|----------|------------------|-------------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
| 1125710 | 300 | Chen | Michael | 42000.00 | 300 | Marketing |
| 1121049 | 400 | Hartson | Robert | 34000.00 | 200 | Research and Dev. |
| 1140131 | 400 | Smith | William | 36000.00 | 300 | Sales |
| 2182221 | 300 | Lakey | Linnea | 42000.00 | 400 | Customer Support |
| 2000000 | 10 | Smith | Steven | 42000.00 | 300 | Human Resources |
| 1100174 | 10 | Smith | Richard | 51000.00 | 300 | Marketing |
| 1121034 | 400 | Whiting | Chris | 34000.00 | 300 | Marketing |
| 1121047 | 300 | Chen | Michael | 42000.00 | 300 | Marketing |
| 1121054 | 200 | Smith | John | 48000.00 | 300 | Marketing |

| LEFT OUTER JOIN | | | First_Name | | Department_Name | |
|-----------------|-----------------------|-----------------|------------|------------------|-----------------|------------------|
| SELECT | First_Name, | Department_Name | Marketing | | Marketing | |
| FROM | Employee_Table as E | Marketing | | Customer Support | | Customer Support |
| LEFT OUTER JOIN | Department_Table as D | Sales | | Sales | | Sales |
| ON | E.Dept_No = D.Dept_No | Richard | | Customer Support | | Customer Support |
| | | Billy | | Research and Dev | | Research and Dev |
| | | John | | Marketing | | Marketing |

A LEFT Outer Join Returns all rows from the LEFT Table, including all Matches. If a LEFT row can't find a match, a NULL is placed on right columns not found!

RIGHT OUTER JOIN

| Employee_Table | | | | | Department_Table | |
|----------------|---------|-----------|------------|----------|------------------|-------------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
| 1125710 | 300 | Chen | Michael | 42000.00 | 300 | Marketing |
| 1121049 | 400 | Hartson | Robert | 34000.00 | 200 | Research and Dev. |
| 1140131 | 400 | Smith | William | 36000.00 | 300 | Sales |
| 2182221 | 300 | Lakey | Linnea | 42000.00 | 400 | Customer Support |
| 2000000 | 10 | Smith | Steven | 42000.00 | 300 | Human Resources |
| 1100174 | 10 | Smith | Richard | 51000.00 | 300 | Marketing |
| 1121034 | 400 | Whiting | Chris | 34000.00 | 300 | Marketing |
| 1121047 | 300 | Chen | Michael | 42000.00 | 300 | Marketing |
| 1121054 | 200 | Smith | John | 48000.00 | 300 | Marketing |

| RIGHT OUTER JOIN | | | First_Name | | Department_Name | |
|------------------|-----------------------|-----------------|------------|------------------|-----------------|------------------|
| SELECT | First_Name, | Department_Name | Marketing | | Marketing | |
| FROM | Employee_Table as E | Marketing | | Customer Support | | Customer Support |
| RIGHT OUTER JOIN | Department_Table as D | Sales | | Sales | | Sales |
| ON | E.Dept_No = D.Dept_No | Richard | | Customer Support | | Customer Support |
| | | Billy | | Research and Dev | | Research and Dev |
| | | John | | Marketing | | Marketing |

This is a RIGHT OUTER JOIN. That means that all rows from the RIGHT Table will appear in the report regardless if it finds a match with the LEFT Table.

RIGHT OUTER JOIN Brings Back All Rows in the RIGHT Table

| Employee_Table | | | | | Department_Table | |
|----------------|---------|------------|------------|---------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
| 1212378 | 100 | Chambers | Madeline | 4880.00 | 100 | Marketing |
| 1296149 | 400 | Hammes | Robert | 5450.00 | 200 | Research and Dev |
| 1242124 | 400 | Billy | William | 3600.00 | 300 | Sales |
| 212225 | 200 | Larkin | Lawrence | 4200.00 | 400 | Customer Support |
| 2000000 | 1 | Jones | Richard | 1200.00 | 500 | Human Resources |
| 1000234 | 10 | Trinkle | Robert | 1400.00 | | |
| 1212324 | 400 | Strickling | Cletus | 5450.00 | | |
| 1214447 | 200 | Coffey | Billy | 4180.00 | | |
| 1333454 | 200 | Smith | John | 4000.00 | | |

| RIGHT OUTER JOIN | | First_Name | Department_Name |
|------------------|-----------------------------|------------|------------------|
| SELECT | First_Name, Department_Name | Madeline | Marketing |
| FROM | Employee_Table as E | Robert | Customer Support |
| RIGHT OUTER JOIN | Department_Table as D | William | Customer Support |
| ON | E.Dept_No = D.Dept_No | Lawrence | Sales |
| | | Cletus | Customer Support |
| | | Billy | Research and Dev |
| | | John | Research and Dev |

All rows from the Right Table were returned with matches and Dept_No 500 didn't have a match, so the system put a NULL Value for Left Column values

FULL OUTER JOIN

| Employee_Table | | | | | Department_Table | |
|----------------|---------|------------|------------|---------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
| 1212378 | 100 | Chambers | Madeline | 4880.00 | 100 | Marketing |
| 1296149 | 400 | Hammes | Robert | 5450.00 | 200 | Research and Dev |
| 1242124 | 400 | Billy | William | 3600.00 | 300 | Sales |
| 212225 | 200 | Larkin | Lawrence | 4200.00 | 400 | Customer Support |
| 2000000 | 1 | Jones | Richard | 1200.00 | 500 | Human Resources |
| 1000234 | 10 | Trinkle | Robert | 1400.00 | | |
| 1212324 | 400 | Strickling | Cletus | 5450.00 | | |
| 1214447 | 200 | Coffey | Billy | 4180.00 | | |
| 1333454 | 200 | Smith | John | 4000.00 | | |

```

SELECT First_Name, Department_Name
FROM Employee_Table as E
FULL OUTER JOIN
  Department_Table as D
ON E.Dept_No = D.Dept_No ;

```

Full Outer will
return all rows
even both tables!

This is a FULL OUTER JOIN. That means that all rows from both the RIGHT and LEFT Table will appear in the report regardless if it finds a match.

FULL OUTER JOIN Brings Back All Rows in All Tables

| Employee_Table | | | | | Department_Table | |
|----------------|---------|-----------|------------|--------|------------------|-----------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |

| | | | | | | |
|---------|-----|-----------|---------|----------|-----|------------------|
| 1231278 | 200 | Chambers | Rebecca | 48890.00 | 100 | Marketing |
| 1236249 | 400 | Harrison | Rebecca | 34500.00 | 200 | Research and Dev |
| 1242115 | 400 | Reilly | William | 30700.00 | 300 | Sales |
| 1212225 | 300 | Larkins | Suzanne | 49200.00 | 400 | Customer Support |
| 1232400 | 7 | Shaw | Gregory | 12800.00 | 500 | Human Resources |
| 1202234 | 10 | Depta | Richard | 12500.00 | | |
| 1212134 | 400 | Strawling | Cletus | 34500.00 | | |
| 1214457 | 200 | Coffing | Billy | 41888.88 | | |
| 1233454 | 200 | Smith | John | 43200.00 | | |

| | | First_Name | Department_Name |
|------------------------------------|--|------------|------------------|
| FULL OUTER JOIN | | | |
| SELECT First_Name, Department_Name | | Rebecca | Marketing |
| FROM Emp_Conte_Table as E | | William | Customer Support |
| FULL OUTER JOIN | | John | Customer Support |
| Department_Table as D | | Richard | Human Resources |
| ON E.Dept_No = D.Dept_No | | Cletus | Customer Support |
| | | Billy | Research and Dev |
| | | John | Research and Dev |

The FULL Outer Join Returns all rows from both Tables. NULLs show the Rows!

Which Tables are the Left and which are the Right?

```

SELECT Cla.Claim_Id,
       Cla.Claim_Date,
       SUB.Last_Name,
       SUB.First_Name,
       "ADD".Phone,
       SER.Service_Pay,
       PRO.Provider_Code,
       PRO.Provider_Name,
FROM CLAIMS Cla
LEFT OUTER JOIN PROVIDERS PRO
  ON Cla.Provider_No = PRO.Provider_Code
LEFT OUTER JOIN SERVICES SER
  ON Cla.Claim_Service = SER.Service_Code
LEFT OUTER JOIN SUBSCRIBERS SUB
  ON Cla.Subscriber_No = SUB.Subscriber_No
AND Cla.Member_No = SUB.Member_No
LEFT OUTER JOIN ADDRESSES "ADD"
  ON SUB.Subscriber_No = "ADD".Subscriber_No

```

Your mission is to show which tables are Left Tables and which ones are Right Tables.

Answer - Which Tables are the Left and which are the Right?


```

SELECT Cla.Claim_Id,
       Cla.Claim_Date,
       SUB.Last_Name,
       SUB.First_Name,
       "ADD".Phone,
       SER.Service_Prov,
       PRO.Provider_Code,
       PRO.Provider_Name
FROM CLAIMS Cla
LEFT OUTER JOIN PROVIDERS PRO
  ON Cla.Provider_No = PRO.Provider_Code
LEFT OUTER JOIN SERVICES SER
  ON Cla.Claim_Service = SER.Service_Code
LEFT OUTER JOIN SUBSCRIBERS SUB
  ON Cla.Subscriber_No = SUB.Subscriber_No
  AND Cla.Member_No = SUB.Member_No
LEFT OUTER JOIN ADDRESSES "ADD"
  ON SUB.Subscriber_No = "ADD".Subscriber_No;

```

There is always
only one left table
(the first table)

All tables after the
first table are each
Right Tables.

The Spool from each
join must contain the
Left Table

The first table is always the left table and the rest are the right tables. It is the spool or intermediate results from each join that remain the left table. In this case all rows will return from the Claims table.

INNER JOIN with Additional AND Clause

| Employee_Table | | | | Department_Table | |
|----------------|---------|------------|----------|------------------|------------------|
| Employee_No | Dept_No | First_Name | Salary | Dept_No | Department_Name |
| 1122170 | 100 | Chambers | 41850.00 | 200 | Marketing |
| 1154149 | 400 | Harrison | 54500.00 | 200 | Research and Dev |
| 1164110 | 400 | Pavilly | 34000.00 | 300 | Sales |
| 2112225 | 300 | Sarkis | 40200.00 | 400 | Customer Support |
| 2000000 | 7 | Jones | 32000.00 | 500 | Human Resources |
| 1000214 | 10 | DePina | 32000.00 | | |
| 1121124 | 400 | Strickling | 54500.00 | | |
| 1124457 | 200 | Coffing | 41850.00 | | |
| 1131074 | 200 | Smith | 41850.00 | | |

```

SELECT First_Name, Last_Name, Department_Name
FROM Employee_Table as E,
     Department_Table as D
WHERE E.Dept_No = D.Dept_No
  AND Department_Name like 'Marketing';

```

The additional AND is performed first in order to eliminate unwanted data, so the join is less intensive than joining everything first and then eliminating.

ANSI INNER JOIN with Additional AND Clause

| Employee_Table | | | | Department_Table | |
|----------------|---------|------------|----------|------------------|------------------|
| Employee_No | Dept_No | First_Name | Salary | Dept_No | Department_Name |
| 1122170 | 100 | Chambers | 41850.00 | 200 | Marketing |
| 1154149 | 400 | Harrison | 54500.00 | 200 | Research and Dev |
| 1164110 | 400 | Pavilly | 34000.00 | 300 | Sales |
| 2112225 | 300 | Sarkis | 40200.00 | 400 | Customer Support |
| 2000000 | 7 | Jones | 32000.00 | 500 | Human Resources |

| Employee_No | Dept_No | First_Name | Last_Name | Salary | Dept_Name |
|-------------|---------|------------|-----------|----------|-----------|
| 1000234 | 10 | Wayne | Richard | 32000.00 | |
| 1121334 | 400 | Prickling | Cletus | 54500.00 | |
| 1324457 | 200 | Coffing | Billy | 41888.88 | |
| 1333454 | 200 | Smith | John | 49000.00 | |

```

SELECT First_Name, Last_Name, Department_Name
FROM Employee_Table as E
INNER JOIN
    Department_Table as D
ON E.Dept_No = D.Dept_No
AND Department_Name like 'Marke%';

```

The additional AND is performed first in order to eliminate unwanted data, so the join is less intensive than joining everything first and then eliminating after.

ANSI INNER JOIN with Additional WHERE Clause

| Employee_No | Dept_No | First_Name | Last_Name | Salary | Dept_Name |
|-------------|---------|------------|-----------|----------|-----------|
| 1000234 | 10 | Wayne | Richard | 32000.00 | |
| 1121334 | 400 | Prickling | Cletus | 54500.00 | |
| 1324457 | 200 | Coffing | Billy | 41888.88 | |
| 1333454 | 200 | Smith | John | 49000.00 | |

```

SELECT First_Name, Last_Name, Department_Name
FROM Employee_Table as E
ANSI INNER JOIN
    Department_Table as D
ON E.Dept_No = D.Dept_No
AND Department_Name like 'Marke%';

```

The additional AND is performed first in order to eliminate unwanted data, so the join is less intensive than joining everything first and then eliminating after.

OUTER JOIN with Additional WHERE Clause

| Employee_No | Dept_No | First_Name | Last_Name | Salary | Dept_Name |
|-------------|---------|------------|-----------|----------|-----------|
| 1000234 | 10 | Wayne | Richard | 32000.00 | |
| 1121334 | 400 | Prickling | Cletus | 54500.00 | |
| 1324457 | 200 | Coffing | Billy | 41888.88 | |
| 1333454 | 200 | Smith | John | 49000.00 | |

The additional WHERE is always performed last on Outer Joins. All rows will be joined first, and then the additional WHERE clause filters after the join takes place.

OUTER JOIN with Additional AND Clause

| Employee Table | | | | | Department Table | |
|----------------|---------|------------|------------|---------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
| 11213750 | 105 | Chandrase | Ramjee | 8950.00 | 100 | Marketing |
| 11212480 | 105 | Chandrase | Ramjee | 5450.00 | 250 | Research and Dev |
| 21042118 | 400 | Malley | William | 3600.00 | 300 | Sales |
| 21042212 | 300 | LeRoux | Lorraine | 6020.00 | 400 | Customer Support |
| 20000000 | 300 | James | James | 5200.00 | 500 | Human Resource |
| 19999994 | 10 | Seymour | Richard | 3200.00 | | |
| 11211394 | 400 | Strickling | Cletus | 5450.00 | | |
| 11214687 | 10 | Coffing | Malley | 4100.00 | | |
| 19999454 | 200 | Smith | John | 4050.00 | | |

```
SELECT First_Name, Department_Name  
FROM Employee_Table as E  
LEFT OUTER JOIN  
    Department_Table as D  
ON E.Dept_No = D.Dept_No  
AND E.Dept_No = 100;
```

The additional AND is performed in conjunction with the ON statement on Outer Joins. All rows will be evaluated with the ON clause and the AND combined.

Results from OUTER JOIN with Additional AND Clause

| Employee Table | | | | Department Table | |
|----------------|-----------|------------|---------|------------------|-------------------|
| Employee_ID | Last_Name | First_Name | Salary | Dept_ID | Department_Name |
| 12345678 | Chambers | Mandaw | 4855.00 | 100 | Marketing |
| 12345679 | Wheat | Marissa | 3450.00 | 100 | Marketing and Ser |
| 12345680 | Bellamy | William | 3600.00 | 100 | Sales |
| 12345681 | Leakins | Lorena | 4020.00 | 400 | Customer Support |
| 12345682 | Johnson | Patricia | 3240.00 | 500 | Human Resources |
| 1200234 | 8 | Smythe | Richard | 3200.00 | |
| 1212394 | 400 | Strickland | Cecilia | 5650.00 | |
| 1204657 | 500 | Smith | William | 4180.00 | |
| 1239456 | 200 | Smith | John | 4800.00 | |

| | First_Name | Department_Name |
|------------------------------------|------------|-----------------|
| OUTER Join with additional AND | Randee | Marketing |
| | Herbert | ? |
| SELECT First_Name, Department_Name | William | ? |
| FROM Employee Table as E | Loraine | ? |

```

LEFT OUTER JOIN
  Department_Table as D
ON E.Dept_No = D.Dept_No
AND E.Dept_No < 100;

```

The additional AND is performed in conjunction with the ON statement on Outer Joins. This can surprise you. Only Mando is in Dept. No. 100!

Quiz - Why is this considered an INNER JOIN?

| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
|-------------|---------|-------------|------------|----------|---------|------------------|
| 1122570 | 400 | Chandrase | Mande | 49500.00 | 200 | Marketing |
| 1124149 | 400 | Martinez | Robert | 54500.00 | 200 | Research and Dev |
| 2341218 | 400 | Pelley | William | 34000.00 | 300 | Sales |
| 2121205 | 300 | Sachdev | Sachdev | 40000.00 | 400 | Customer Support |
| 2000000 | ? | Jones | Spikey | 32000.00 | 500 | Human Resources |
| 2000294 | 10 | Byrnes | Richard | 32000.00 | | |
| 1121334 | 400 | Potterkling | Celia | 54500.00 | | |
| 1324487 | 200 | Coffing | Billy | 41880.00 | | |
| 1324854 | 200 | Smith | John | 49000.00 | | |

```

SELECT First_Name, Department_Name
FROM Employee_Table as E
LEFT OUTER JOIN
  Department_Table as D
ON E.Dept_No = D.Dept_No
AND Department_Name like 'Marketing';

```

This is considered an INNER JOIN because we are doing a LEFT OUTER JOIN on the Employee_Table, and then filtering with the AND for a column in the right table!

The DREADED Product Join

| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
|-------------|---------|-------------|------------|----------|---------|------------------|
| 1122570 | 400 | Chandrase | Mande | 49500.00 | 200 | Marketing |
| 1124149 | 400 | Martinez | Robert | 54500.00 | 200 | Research and Dev |
| 2341218 | 400 | Pelley | William | 34000.00 | 300 | Sales |
| 2121205 | 300 | Sachdev | Sachdev | 40000.00 | 400 | Customer Support |
| 2000000 | ? | Jones | Spikey | 32000.00 | 500 | Human Resources |
| 2000294 | 10 | Byrnes | Richard | 32000.00 | | |
| 1121334 | 400 | Potterkling | Celia | 54500.00 | | |
| 1324487 | 200 | Coffing | Billy | 41880.00 | | |
| 1324854 | 200 | Smith | John | 49000.00 | | |

```

SELECT First_Name, Last_Name, Department_Name
FROM Employee_Table as E
WHERE Department_Table as D
WHERE Department_Name like 'Marketing';
Order by 1, 2, 3;

```

No Join
Condition
Linking the
Two Tables!

This query becomes a Product Join because it does not possess any JOIN Conditions (Join Keys). Every row from one table is compared to every row of the other table, and quite often the data is not what you intended to get back.

Result Set of the DREADED Product Join

No Join Condition, Linking the Two Tables

No null returned here

```
SELECT First_Name, Last_Name, Department_Name
FROM Employee_Table as E,
Department_Table as D
WHERE Department_Name like 'Sales';
Order by 1, 2, 3;
```

| First_Name | Last_Name | Department_Name |
|------------|-----------|------------------|
| Billy | Coffey | Customer Support |
| Billy | Coffey | Human Resources |
| Billy | Coffey | Marketing |
| Chen | Wickling | Customer Support |
| Chen | Wickling | Human Resources |
| Chen | Wickling | Marketing |
| DeHaan | Herman | Customer Support |
| DeHaan | Herman | Human Resources |
| DeHaan | Herman | Marketing |

27 Rows cause back 9 employees with each working in 3 different departments simultaneously? Duplicate?

How can Billy Coffey work in 3 different departments?

A Product Join is often a mistake! Three Department rows had an 'm' in their name. These were joined to every employee and the information is worthless.

The Horrifying Cartesian Product Join

| Employee_Id | Dept_Id | First_Name | Last_Name | Salary | Dept_Id | Department_Name |
|-------------|---------|------------|-----------|----------|---------|-------------------|
| 1212878 | 200 | Chambeca | Handee | 40550.00 | 200 | Marketing |
| 1212880 | 200 | Barclison | Rebecca | 40550.00 | 200 | Marketing and Dev |
| 2341111 | 200 | Bailey | William | 34000.00 | 200 | Sales |
| 2311123 | 200 | Jankins | Loislane | 42000.00 | 200 | Customer Support |
| 2300100 | 2 | Shane | Equipp | 22000.00 | 200 | Human Resources |
| 1000294 | 10 | Myrland | Richard | 22000.00 | 200 | Human Resources |
| 1211114 | 400 | Stefanling | Carson | 44000.00 | 200 | Human Resources |
| 1214057 | 200 | Coffey | Billy | 41000.00 | 200 | Human Resources |
| 1214054 | 200 | Smith | John | 40000.00 | 200 | Human Resources |

No WHERE Clause in the join!

```
SELECT First_Name, Last_Name, Department_Name
FROM Employee_Table as E,
Department_Table as D
```

This pairs every row from one table to every row of another table. Nine rows multiplied by 5 rows = 45 rows of complete nonsense!

A Cartesian Product Join is a big mistake.

The ANSI Cartesian Join will ERROR

| Employee_Table | Department_Table |
|----------------|------------------|
| | |

| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
|-------------|---------|------------|------------|----------|---------|------------------|
| 1123570 | 400 | Chambers | Heather | 49500.00 | 400 | Marketing |
| 1124165 | 400 | Harrison | Hebert | 54500.00 | 300 | Research and Dev |
| 2342218 | 400 | Beilly | William | 36000.00 | 300 | Sales |
| 2112225 | 300 | Larkins | Leanne | 40200.00 | 400 | Customer Support |
| 2000000 | ? | Jones | Elvis | 32000.00 | 500 | Human Resources |
| 1000234 | 10 | Seytze | Richard | 32000.00 | | |
| 1121334 | 400 | Strickling | Clatus | 54500.00 | | |
| 1124687 | 200 | Coffey | Billy | 41888.88 | | |
| 1333454 | 200 | Smith | John | 49000.00 | | |

```
SELECT First_Name, Last_Name, Department_Name
FROM Employee_Table as E
INNER JOIN
Department_Table as D
```

No ON
Clause in
the join!

This causes an error. ANSI won't let this run unless a join condition is present.

Quiz - Do these Joins Return the Same Answer Set?

| Employee_Table | | | | | Department_Table | |
|----------------|---------|------------|------------|----------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
| 1123570 | 400 | Chambers | Heather | 49500.00 | 400 | Marketing |
| 1124165 | 400 | Harrison | Hebert | 54500.00 | 300 | Research and Dev |
| 2342218 | 400 | Beilly | William | 36000.00 | 300 | Sales |
| 2112225 | 300 | Larkins | Suzanne | 40200.00 | 400 | Customer Support |
| 2000000 | ? | Jones | Elvis | 32000.00 | 500 | Human Resources |
| 1000234 | 10 | Seytze | Richard | 32000.00 | | |
| 1121334 | 400 | Strickling | Clatus | 54500.00 | | |
| 1124687 | 200 | Coffey | Billy | 41888.88 | | |
| 1333454 | 200 | Smith | John | 49000.00 | | |

| | | | | | | |
|------------------------------------|--|--|--|--|------------------------------------|--|
| /* Query 1 */ | | | | | /* Query 2 */ | |
| SELECT First_Name, Department_Name | | | | | SELECT First_Name, Department_Name | |
| FROM Employee_Table | | | | | FROM Employee_Table, | |
| INNER JOIN | | | | | Department_Table | |
| | | | | | Department_Table | |

Do these two queries produce the same result?

Answer - Do these Joins Return the Same Answer Set?

| Employee_Table | | | | | Department_Table | |
|----------------|---------|------------|------------|----------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
| 1123570 | 400 | Chambers | Heather | 49500.00 | 400 | Marketing |
| 1124165 | 400 | Harrison | Hebert | 54500.00 | 300 | Research and Dev |
| 2342218 | 400 | Beilly | William | 36000.00 | 300 | Sales |
| 2112225 | 300 | Larkins | Leanne | 40200.00 | 400 | Customer Support |
| 2000000 | ? | Jones | Elvis | 32000.00 | 500 | Human Resources |
| 1000234 | 10 | Seytze | Richard | 32000.00 | | |
| 1121334 | 400 | Strickling | Clatus | 54500.00 | | |
| 1124687 | 200 | Coffey | Billy | 41888.88 | | |
| 1333454 | 200 | Smith | John | 49000.00 | | |

| | | | | | | |
|---------------|--|--|--|--|---------------|--|
| /* Query 1 */ | | | | | /* Query 2 */ | |
|---------------|--|--|--|--|---------------|--|

```

SELECT First_Name,
       Department_Name
FROM Emp_Layer_Table
INNER JOIN
       Department_Table

```

```

SELECT First_Name,
       Department_Name
FROM Emp_Layer_Table
FROM Department_Table

```

Do these two queries produce the same result? No, Query 1 Errors due to ANSI syntax and no ON Clause, but Query 2 Product Joins to bring back junk!

The CROSS JOIN

| Customer_Table | | Order_Table | |
|-----------------|---------------------|--------------|-------------|
| Customer_Number | Customer_Name | Order_Number | Order_Total |
| 01111111 | Billy's Best Choice | 123456 | 11111111 |
| 01121111 | Acme Products | 123512 | 11111111 |
| 01121124 | ACT Consulting | 123562 | 11121124 |
| 01940803 | XYZ Plumbing | 123585 | 07122454 |
| 07121456 | Databases N-L | 123777 | 07094693 |

A Cross Join
is the ANSI
equivalent to
a Product Join

```

SELECT Customer_Name,
       Order_Number
FROM Customer_Table
CROSS JOIN
       Order_Table
WHERE Order_Number = 123456
ORDER BY 1;

```

Only a WHERE
will work. ON
Will NOT!

This query becomes a Product Join because a Cross Join is an ANSI Product Join. It will compare every row from the Customer_Table to Order_Number 123456 in the Order_Table. Check out the Answer Set on the next page.

The CROSS JOIN Answer Set

| Customer_Table | | Order_Table | |
|-----------------|---------------------|--------------|-------------|
| Customer_Number | Customer_Name | Order_Number | Order_Total |
| 11111111 | Billy's Best Choice | 123456 | 11111111 |
| 31313131 | Acme Products | 123512 | 11111111 |
| 31323134 | ACT Consulting | 123552 | 31323134 |
| 97668083 | XYZ Plumbing | 123585 | 5111.47 |
| 87323456 | Databases N-U | 123777 | 87323456 |
| | | | 15231.62 |
| | | | 23454.14 |

```

SELECT Customer_Name,
       Order_Number
FROM Customer_Table
CROSS JOIN
       Order_Table
WHERE Order_Number = 123456
ORDER BY 1;

```

Answer
Set

| Customer_Name | Order_Number |
|---------------------|--------------|
| Billy's Best Choice | 123456 |
| Acme Products | 123456 |
| ACT Consulting | 123456 |
| XYZ Plumbing | 123456 |
| Databases N-L | 123456 |

This Cross Join produces information that quite often just isn't worth anything!

The Self Join

| Employee_Table2 | | | | | |
|-----------------|---------|-----------|------------|----------|-----|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Mgr |
| 1222779 | 100 | Chaudhry | Shamim | 48350.00 | Y |
| 1256149 | 400 | Harrison | Herbert | 54500.00 | N |
| 2341218 | 400 | Reilly | William | 36000.00 | Y |
| 1121344 | 400 | Stevens | Chris | 54500.00 | N |
| 2312225 | 300 | Lodman | Louise | 60200.00 | Y |
| 2050060 | 7 | Tanev | Svetlana | 32000.50 | N |
| 1006214 | 10 | Suytnae | Richard | 32800.00 | N |
| 1124957 | 200 | Coffing | Billy | 41888.88 | N |
| 1313551 | 200 | Suans | John | 48000.00 | Y |

Notice we
Added a
column
which
will tell us
if someone
is a
manager

```
SELECT MgrsDept_No, MgrsLast_Name as MgrName, MgrsSalary as MgrSal,
EmpsLast_Name as EmpName, EmpsSalary as EmpSal
FROM Employee_Table2 as Emps,
Employee_Table2 as Mgrs
WHERE Emps.Dept_No = Mgrs.Dept_No
AND Emps.Mgr = 'Y'
```

Which Workers
make a bigger
Salary than
their Manager?

A Self Join gives itself 2 different Aliases, which is then seen as two different tables.

The Self Join with ANSI Syntax

| Employee_Table2 | | | | | |
|-----------------|---------|-----------|------------|----------|-----|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Mgr |
| 122276 | 110 | Cheney | Mindee | 6188.00 | Y |
| 125049 | 400 | Harmon | Buckett | 54500.00 | N |
| 2541216 | 400 | Reilly | William | 30000.00 | Y |
| 1161164 | 400 | Stuchling | Chen | 54500.00 | N |
| 23112125 | 300 | Luskus | Loonae | 42200.00 | Y |
| 2005000 | 7 | Stane | Spigge | 32800.50 | N |
| 1009244 | 10 | Stuyler | Richard | 32800.00 | N |
| 1324857 | 200 | Coffing | Billy | 41888.88 | N |
| 1321454 | 200 | Smith | Jake | 45000.00 | Y |

Notice we added a column which will tell us if someone is a manager

```
SELECT Mgrs.Dept_No, Mgrs.Last_Name as MgrsName, Mgrs.Salary as MgrsSal,
Emps.Last_Name as EmpsName, Emps.Salary as EmpsSal
FROM Employee_Table2 as Emps
INNER JOIN
Employee_Table2 as Mgrs
ON Emps.Dept_No = Mgrs.Dept_No
WHERE Mgrs.Mgr = 'Y'
AND Emps.Salary > Mgrs.Salary;
```

Which Workers make a bigger Salary than their Manager?

A Self Join gives itself 2 different Aliases, which is then seen as two different tables.

Quiz - Will both queries bring back the same Answer Set?

| Customer_Table | | Order_Table | |
|-----------------|--------------------|--------------|-------------|
| Customer_Number | Customer_Name | Order_Number | Order_Total |
| 01111111 | Bill's Best Choice | 123456 | 12345.67 |
| 01111111 | Acme Products | 123512 | 8009.91 |

| | | | |
|-----------------------------|--------|----------|----------|
| 3123134 ACME Communications | 123552 | 3123134 | 1111.47 |
| 97094893 XYZ Plumbing | 123553 | 97224454 | 12351.42 |
| 97024454 Database D-C | 123777 | 97094893 | 23456.58 |

| | |
|---|---|
| SELECT Customer_Name, Order_Number, Order_Total FROM Customer_Table as Cust ORDER BY 2018 Order_Table as ORD ON Cust.Customer_Number = Ord.Customer_Number WHERE Customer_Name like 'Billy' | SELECT Customer_Name, Order_Number, Order_Total FROM Customer_Table as Cust ORDER BY 2018 Order_Table as ORD ON Cust.Customer_Number = Ord.Customer_Number AND Customer_Name like 'Billy' |
|---|---|

Will both queries bring back the same result set?

Answer – Will both queries bring back the same Answer Set?

| Customer_Table | Customer_Name | Order_Number | Order_Table | Customer_Number | Order_Total |
|----------------|---------------------|--------------|-------------|-----------------|-------------|
| 3123134 | Billy's Best Choice | 123456 | 3123134 | 1111111 | 12347.53 |
| 3123131 | Acme Products | 123512 | 3123131 | 1111111 | 8005.91 |
| 3123134 | ACME Communications | 123512 | 3123134 | 1111.47 | |
| 97094893 | XYZ Plumbing | 123553 | 97224454 | 12351.42 | |
| 97024454 | Database D-C | 123777 | 97094893 | 23456.58 | |

| | |
|---|---|
| SELECT Customer_Name, Order_Number, Order_Total FROM Customer_Table as Cust ORDER BY 2018 Order_Table as ORD ON Cust.Customer_Number = Ord.Customer_Number WHERE Customer_Name like 'Billy' | SELECT Customer_Name, Order_Number, Order_Total FROM Customer_Table as Cust ORDER BY 2018 Order_Table as ORD ON Cust.Customer_Number = Ord.Customer_Number AND Customer_Name like 'Billy' |
|---|---|

Will both queries bring back the same result set? Yes! They are both inner joins, so they bring back the same results. Had they been outer joins, the WHERE and the AND would act differently.

Quiz – Will both queries bring back the same Answer Set?

| Customer_Table | Customer_Name | Order_Number | Order_Table | Customer_Number | Order_Total |
|----------------|---------------------|--------------|-------------|-----------------|-------------|
| 3123134 | Billy's Best Choice | 123456 | 3123134 | 1111111 | 12347.53 |
| 3123131 | Acme Products | 123512 | 3123131 | 1111111 | 8005.91 |
| 3123134 | ACME Communications | 123512 | 3123134 | 1111.47 | |
| 97094893 | XYZ Plumbing | 123553 | 97224454 | 12351.42 | |
| 97024454 | Database D-C | 123777 | 97094893 | 23456.58 | |

| | |
|--|--|
| SELECT Customer_Name, Order_Number, Order_Total FROM Customer_Table as Cust LEFT OUTER JOIN ORDER BY 2018 Order_Table as ORD ON Cust.Customer_Number = Ord.Customer_Number WHERE Customer_Name like 'Billy' | SELECT Customer_Name, Order_Number, Order_Total FROM Customer_Table as Cust LEFT OUTER JOIN ORDER BY 2018 Order_Table as ORD ON Cust.Customer_Number = Ord.Customer_Number AND Customer_Name like 'Billy' |
|--|--|

Will both queries bring back the same result set?

Answer – Will both queries bring back the same Answer Set?

| Customer_Table | | Order_Table | |
|-----------------|---------------------|--------------|-------------------|
| Customer_Number | Customer_Name | Order_Number | Order_Total |
| 11111111 | Billy's Best Choice | 123456 | 11111111 12345.53 |
| 11111111 | Rome Furniture | 123512 | 11111111 8005.74 |
| 11223344 | ACE Consulting | 123562 | 11223344 5111.47 |
| 57894321 | RTZ Plumbing | 123585 | 57894321 15231.42 |
| 57321456 | Databassz B-C | 123777 | 57321456 23456.14 |

```
SELECT Customer_Name,
       Order_Number,
       Order_Total
FROM Customer_Table as Cust
LEFT OUTER JOIN
       Order_Table as ORD
ON Cust.Customer_Number
= Ord.Customer_Number
WHERE
       Customer_Name like 'Billy*'
ORDER BY 1;
```

```
SELECT Customer_Name,
       Order_Number,
       Order_Total
FROM Customer_Table as Cust
LEFT OUTER JOIN
       Order_Table as ORD
ON Cust.Customer_Number
= Ord.Customer_Number
AND
       Customer_Name like 'Billy*'
ORDER BY 1;
```

Will both queries bring back the same result set? NO! The WHERE is performed last.

How would you Join these two tables?

| Course_Table | | | |
|--------------|--------------------------|---------|-------|
| Course_ID | Course_Name | Credits | Seats |
| 100 | Database Concepts | 3 | 50 |
| 200 | Introduction to SQL | 3 | 20 |
| 210 | Advanced SQL | 3 | 20 |
| 220 | Virtual SQL Features | 2 | 20 |
| 300 | Physical Database Design | 4 | 20 |
| 400 | Database Administration | 4 | 16 |

| Student_Table | | | | |
|---------------|-----------|------------|------------|----------|
| Student_ID | Last_Name | First_Name | Class_Code | Grade_Pt |
| 423400 | Larkins | Michael | FR | 0.00 |
| 231222 | Wilkins | Suzie | SO | 3.00 |
| 280023 | McRoberts | Richard | JR | 1.90 |
| 321133 | Bond | Jimmy | JR | 3.95 |
| 123624 | Samson | Henry | FR | 2.85 |
| 333450 | Smith | Andy | SO | 2.00 |
| 324832 | Delaney | Danny | JR | 3.35 |
| 260000 | Johnson | Stanley | ? | ? |
| 234121 | Thomas | Wendy | FR | 4.00 |
| 192350 | Phillips | Marvin | SR | 3.00 |

Looking at the columns above, which will allow these two tables to join?

How would you join these two tables? You Can't Yet!

| Course_Table | | | |
|--------------|---------------------|---------|-------|
| Course_ID | Course_Name | Credits | Seats |
| 100 | Database Concepts | 3 | 50 |
| 200 | Introduction to SQL | 3 | 20 |
| 210 | Advanced SQL | 3 | 20 |

| | | | |
|-----|--------------------------|---|----|
| 220 | VB3 SQL Features | 2 | 20 |
| 300 | Physical Database Design | 4 | 20 |
| 400 | Database Administration | 4 | 16 |

| Student_ID | Last_Name | First_Name | Class_Code | Grade_Pt |
|------------|-----------|------------|------------|----------|
| 4029402 | Larkins | Michael | FR | 0.00 |
| 2512721 | William | Susan | SO | 3.00 |
| 2000223 | McRoberts | Richard | JR | 1.00 |
| 3221131 | Bond | Jimmy | JR | 1.00 |
| 1155434 | Ransom | Henry | FR | 2.00 |
| 3118450 | Smith | Andy | SO | 1.00 |
| 3246552 | Delaney | Danny | SR | 3.35 |
| 2400000 | Johnson | Stanley | F | 1 |
| 2141121 | Thomas | Wendy | FR | 4.00 |
| 1212552 | Phillips | Martin | FR | 3.00 |

Get ready for the Associative Table!

An Associative Table is a Bridge that Joins Two Tables

| Associative Table | | Course_Table | | | |
|-------------------|-----------|--------------|--------------------------|---------|-------|
| Student_ID | Course_ID | Course_ID | Course_Seq | Credits | Grade |
| 2000223 | 210 | 200 | Database Concepts | 3 | 10 |
| 2812222 | 210 | 200 | Introduction to SQL | 3 | 20 |
| 1206254 | 300 | 310 | Advanced SQL | 3 | 20 |
| 2812222 | 320 | 220 | VB3 SQL Features | 2 | 20 |
| 1706304 | 300 | 400 | Physical Database Design | 4 | 20 |
| 3221232 | 220 | 400 | Database Administration | 4 | 16 |
| 1206304 | 320 | | | | |
| 3221232 | 300 | | | | |
| 2140552 | 200 | | | | |
| 1834500 | 300 | | | | |
| 2800000 | 400 | | | | |
| 3834500 | 400 | | | | |
| 2841211 | 100 | | | | |
| 1232550 | 100 | | | | |

The Associative Table is a bridge between the Course_Table and Student_Table.


Quiz – Can you Write the 3-Table Join?

| Associative Table | | Course_Table | | | |
|----------------------|-----|--------------|--------------------------|---------|-------|
| Student_Course_Table | | Course_ID | Course_Table | Credits | Grade |
| 200023 | 210 | 100 | Database Concepts | 3 | 50 |
| 200023 | 210 | 200 | Introduction to SQL | 3 | 20 |
| 200023 | 210 | 210 | Advanced SQL | 3 | 10 |
| 200023 | 210 | 220 | VB3 SQL Features | 2 | 20 |
| 200023 | 210 | 300 | Physical Database Design | 4 | 20 |
| 200023 | 210 | 400 | Database Administration | 4 | 10 |

| Student_Table | | Student_Table | | | |
|---------------|-----------|---------------|-----------|------------|------------|
| Student_ID | Course_ID | Student_ID | Last_Name | First_Name | Class_Role |
| 200023 | 210 | 423400 | Larkins | Richard | PR |
| 200023 | 210 | 232222 | Wilson | James | DR |
| 200023 | 210 | 200023 | McRoberts | Richard | DR |
| 200023 | 210 | 322189 | Wood | James | DR |
| 200023 | 210 | 125434 | Hasson | Henry | PR |
| 200023 | 210 | 322450 | Smith | Andy | DR |
| 200023 | 210 | 324652 | Delaney | Sammy | DR |
| 200023 | 210 | 240100 | Johnson | Frankley | DR |
| 200023 | 210 | 234121 | Thomas | Wendy | PR |
| 200023 | 210 | 123250 | Phillips | Martin | DR |


SELECT ALL Columns from the Course_Table and Student_Table and join them.

Answer to Quiz - Can you Write the 3-Table Join?




The Associative Table is a bridge between the Course_Table and Student_Table, and its sole purpose is to join these two tables together.

Quiz - Can you Write the 3-Table Join to ANSI Syntax?




Please re-write the above query using ANSI Syntax.

Answer – Can you Write the 3-Table Join to ANSI Syntax?




The above query has been written using ANSI Syntax.

Quiz – Can you Place the ON Clauses at the End?




Please re-write the above query and place both ON Clauses at the end.

Answer – Can you Place the ON Clauses at the End?




This is tricky. The only way it works is to place the ON clauses backwards. The first ON Clause represents the last INNER JOIN, and then moves backwards.

The S-Table Join – Logical Insurance Model



Above is the logical model for the insurance tables showing the Primary Key and Foreign Key relationships (PKFK).

Quiz - Write a Five Table Join Using ANSI Syntax



Your mission is to write a five table join selecting all columns using ANSI syntax.


Answer - Write a Five Table Join Using ANSI Syntax

```

SELECT
  clai.*, sub1.*, add1.*, prol.*, ser1.*
FROM
  CLAIMS AS clai
INNER JOIN
  SUBSCRIBERS AS sub1
    ON clai.Subscriber_No = sub1.Subscriber_No
    AND clai.Member_No = sub1.Member_No
INNER JOIN
  ADDRESSES AS add1
    ON sub1.Subscriber_No = add1.Subscriber_No
INNER JOIN
  PROVIDERS AS prol
    ON clai.Provider_No = prol.Provider_Code
INNER JOIN
  SERVICES AS ser1
    ON clai.Claim_Service = ser1.Service_Code ;

```

Above is the example writing this five table join using ANSI syntax.

Quiz - Write a Five Table Join Using ANSI Syntax


Your mission is to write a five table join selecting all columns using ANSI syntax.

Answer - Write a Five Table Join Using ANSI Syntax

```
SELECT
  clal.*, sub1.*, add1.*, prol.*, ser1.*
FROM
  CLAIMS AS clal
INNER JOIN
  SUBSCRIBERS AS sub1
    ON clal.Subscriber_No = sub1.Subscriber_No
    AND clal.Member_No = sub1.Member_No
INNER JOIN
  ADDRESSES AS add1
    ON sub1.Subscriber_No = add1.Subscriber_No
INNER JOIN
  PROVIDERS AS prol
    ON clal.Provider_No = prol.Provider_Code
INNER JOIN
  SERVICES AS ser1
    ON clal.Claim_Service = ser1.Service_Code ;
```

Above is the example writing this five table join using ANSI syntax.

Quiz - Write a Five Table Join Using Non-ANSI Syntax



Your mission is to write a five table join selecting all columns using Non-ANSI syntax.

Answer - Write a Five Table Join Using Non-ANSI Syntax

```
SELECT clai.*, sub1.*, add1.*, prov1.*, serv1.*
FROM   CLAIMS AS clai,
       SUBSCRIBERS AS sub1,
       ADDRESSES AS add1,
       PROVIDERS AS prov1,
       SERVICES AS serv1
WHERE  clai.Subscriber_No = sub1.Subscriber_No
AND    clai.Member_No = sub1.Member_No
AND    sub1.Subscriber_No = add1.Subscriber_No
AND    clai.Provider_No = prov1.Provider_Code
AND    clai.Claim_Service = serv1.Service_Code ;
```

Above is an example of writing this five table join using Non-ANSI syntax.

Quiz –Re-Write this putting the ON clauses at the END

```
SELECT clai.*, sub1.*, add1.*, prol.*, ser1.*
FROM CLAIMS AS clai
INNER JOIN SUBSCRIBERS AS sub1
ON clai.Subscriber_No = sub1.Subscriber_No
AND clai.Member_No = sub1.Member_No
INNER JOIN ADDRESSES AS add1
ON sub1.Subscriber_No = add1.Subscriber_No
INNER JOIN PROVIDERS AS prol
ON clai.Provider_No = prol.Provider_Code
INNER JOIN SERVICES AS ser1
ON clai.Claim_Service = ser1.Service_Code ;
```


Above is an example of writing this five table join using Non-ANSI syntax.

Answer –Re-Write this putting the ON clauses at the END

```
SELECT clai.*, sub1.*, add1.*, prol.*, ser1.*
FROM PROVIDERS AS prol
INNER JOIN ADDRESSES AS add1
ON sub1.Subscriber_No = add1.Subscriber_No
INNER JOIN SUBSCRIBERS AS sub1
ON sub1.Subscriber_No = add1.Subscriber_No
INNER JOIN CLAIMS AS clai
ON clai.Claim_Service = ser1.Service_Code
ON clai.Subscriber_No = sub1.Subscriber_No
AND clai.Member_No = sub1.Member_No
ON sub1.Subscriber_No = add1.Subscriber_No
ON clai.Provider_No = prol.Provider_Code ;
```

Above is an example of writing this five table join using ANSI syntax with the ON clauses at the end. Also to make this happen, we had to move the tables around. Notice that the first ON clause represents the last two tables being joined and then it works backwards.

The Nexus Query Chameleon Writes the SQL for Users.



Join Indexes

Join Indexes

Creating a Multi-Table Join Index

```
CREATE MULTI-INDEX EMP_DEPT_JOIN AS
SELECT Employee_No
      , Emp_No
      , First_Name
      , Last_Name
      , Salary
      , Department_Name
FROM   Employee_Table as E
WHERE  EMP_No = E.Emp_No
ORDER BY EMP_No (Employee_No) ;
```

The Syntax above will create a Multi-Table Join Index with a NUP1 on Employee_No. The next slide will illustrate a visual so you can see the data in the Join Index. Join indexes are created so data doesn't have to move to satisfy the join. The Join Index essentially pre-joins the table and keeps it hidden for the Parsing Engine to utilize.

Visual of a Join Index

| Employee_Table | | | | | | Department_Table | |
|----------------|--------|------------|------------|----------|---------|------------------|--|
| Employee_No | Emp_No | Last_Name | First_Name | Salary | Dept_No | Department_Name | |
| 11212876 | 200 | Chambers | Madew | 48850.00 | 500 | Marketing | |
| 11212818 | 400 | Barrington | Madew | 58800.00 | 200 | Research and Dev | |
| 11212123 | 400 | Reilly | William | 36000.00 | 300 | Sales | |
| 11212256 | 300 | Cartlidge | Michael | 40200.00 | 400 | Customer Support | |
| 11212006 | 1 | Jones | Emily | 32800.00 | 500 | Ruman Resources | |
| 11212334 | 10 | Mythen | Richard | 32800.00 | | | |
| 11212394 | 400 | Strickling | Clatus | 48800.00 | | | |
| 11214457 | 200 | Coffey | Billy | 41888.88 | | | |
| 11212454 | 200 | Smith | John | 48000.00 | | | |

| Join Index named EMP_DEPT_JOIN | | | | | | | |
|--------------------------------|--------|------------|------------|----------|------------------|------------------|--|
| Employee_No | Emp_No | Last_Name | First_Name | Salary | Dept_No | Department_Name | |
| 11212876 | 200 | Chambers | Madew | 48850.00 | 500 | Marketing | |
| 11212818 | 400 | Barrington | Madew | 58800.00 | 200 | Customer Support | |
| 11212123 | 400 | Reilly | William | 36000.00 | 300 | Customer Support | |
| 11212256 | 300 | Cartlidge | Michael | 40200.00 | 400 | Sales | |
| 11212334 | 400 | Strickling | Clatus | 48800.00 | Customer Support | | |
| 11214457 | 200 | Coffey | Billy | 41888.88 | Research and Dev | | |
| 11212454 | 200 | Smith | John | 48000.00 | Research and Dev | | |

The Join Index looks like an Answer Set, but each row is stored like a normal table in that the rows of the Join Index are signed amongst the AMPs. Users can't query the Join Index, but the Parsing Engine gets data from the Join Index when it chooses.

Outer Join Multi-Table Join Index


```

CREATE JOIN INDEX EMP_DEPT_LEFTY AS
SELECT Employee_No,
       E_Dept_No,
       First_Name,
       Last_Name,
       Salary,
       Department_Name
FROM Employee_Table as E
LEFT OUTER JOIN
       Department_Table as D
ON   E_Dept_No = D_Dept_No
PRIMARY INDEX (Employee_No) ;

```

A Multi-Table Outer Join Index has some very specific rules above to remember. Turn the page to see a visual of the data.

Visual of a Left Outer Join Index

| Employee_Table | | | | | Department_Table | |
|----------------|---------|-----------|------------|----------|------------------|------------------|
| Employee_No | Dept_No | Last_Name | First_Name | Salary | Dept_No | Department_Name |
| 221208 | 100 | Chambers | Wendee | 48500.00 | 100 | Marketing |
| 221209 | 400 | Barrows | Robert | 34500.00 | 200 | Research and Dev |
| 294128 | 400 | Reilly | William | 34000.00 | 300 | Sales |
| 291228 | 300 | Larkins | Scotline | 40200.00 | 400 | Customer Support |
| 200000 | ? | Jones | Spiggy | 32000.00 | 500 | Human Resources |
| 100128 | 400 | Byrnes | Richard | 32000.00 | ? | |
| 112138 | 400 | Priskling | Cletus | 54500.00 | 400 | Customer Support |
| 131403 | 200 | Coffey | Billy | 41800.00 | 200 | Research and Dev |
| 133104 | 200 | Smith | John | 49000.00 | 200 | Research and Dev |

The Outer Join Index has the additional rows that did NOT match

Compressed Multi-Table Join Index

```

CREATE COMP INDEX Cust_Order_IDX AS
SELECT
  (C-Customer_Number, Customer_Name),
  (Order_Number, Order_Date, Order_Detail)
FROM   Customer_Table as C
ORDER BY
  Order_Table as O
OR
  (C-Customer_Number,
  PRUNGRD_CUSTS (Customer_Number) ) ;

```

A Compressed Multi-Table Join Index won't keep repeating the same Customer_Number and Customer_Name, but only list it once. A visual example follows on the next page.

A Visual of a Compressed Multi-Table Join Index

| Customer_Table | | | Orders_Table | | |
|-----------------|---------------------|--------------|-----------------|-------------|--|
| Customer_Number | Customer_Name | Order_Number | Customer_Number | Order_Total | |
| 11111111 | Billy's Best Choice | 123456 | 11111111 | 12347.53 | |
| 11111111 | Billy's Best Choice | 123456 | 11111111 | 12347.53 | |
| 11111111 | Billy's Best Choice | 123456 | 11111111 | 12347.53 | |
| 11111111 | Billy's Best Choice | 123456 | 11111111 | 12347.53 | |
| 11111111 | Billy's Best Choice | 123456 | 11111111 | 12347.53 | |
| 11111111 | Billy's Best Choice | 123456 | 11111111 | 12347.53 | |
| 11111111 | Billy's Best Choice | 123456 | 11111111 | 12347.53 | |
| 11111111 | Billy's Best Choice | 123456 | 11111111 | 12347.53 | |
| 11111111 | Billy's Best Choice | 123456 | 11111111 | 12347.53 | |

Not
Repeted

Join Index named CUST_ORDER_IDX

| Customer Number | Customer Name | Order Number | Order Date | Order Total |
|-----------------|---------------------|--------------|------------|-------------|
| 11111111 | Billy's Best Choice | 123456 | 05/04/1999 | 12347.53 |
| 11111111 | Billy's Best Choice | 123456 | 05/04/1999 | 12347.53 |
| 11111111 | Billy's Best Choice | 123456 | 05/04/1999 | 12347.53 |
| 11111111 | Billy's Best Choice | 123456 | 05/04/1999 | 12347.53 |
| 11111111 | Billy's Best Choice | 123456 | 05/04/1999 | 12347.53 |
| 11111111 | Billy's Best Choice | 123456 | 05/04/1999 | 12347.53 |
| 11111111 | Billy's Best Choice | 123456 | 05/04/1999 | 12347.53 |
| 11111111 | Billy's Best Choice | 123456 | 05/04/1999 | 12347.53 |
| 11111111 | Billy's Best Choice | 123456 | 05/04/1999 | 12347.53 |
| 11111111 | Billy's Best Choice | 123456 | 05/04/1999 | 12347.53 |

Billy's Best Choice is Customer Number 11111111 and they have placed two orders, but the Customer Number and Customer Name don't repeat unnecessarily.


Creating a Single-Table Join Index

```
CREATE JOIN INDEX Employee_IDX
AS
SELECT Employee_No,
Dept_No,
First_Name,
Last_Name,
Salary
FROM Employee_Table
PRIMARY INDEX (Dept_No);
```

We've duplicated the Employee_Table with a different Primary Index.

If a USER queries with the Dept_No in the WHERE clause this will be a Single-AMP retrieve. If the USER joins the Employee and Department Tables together then Teradata won't need to Redistribute or Duplicate to get the data AMP local. The next page will give you a visual of how that looks on a particular AMP.

Conceptual of a Single-Table Join Index on an AMP



Notice the Primary Indexes on both tables and the Single Table Join Index. The Join Index gives the Parsing Engine options. If a query is run against the Employee_Table with Employee_No in the WHERE clause, it will use the normal table, but if a user uses Dept_No in the WHERE clause, it will use the Join Index. If a user needs to join the Department_Table to the Employee_Table, the Join Index is used so no data moves.

Single Table Join Index Great For LIKE Clause

Build a STJ with column that contains three columns. They are the LIKE column being queried, the primary index column of the base table and the keyword ROWID!

```
CREATE JOIN INDEX Good_Ford_Like AS
SELECT S1License_Plate#, P1_Col, ROWID/* Global Join Index */
FROM BDF_Table
Primary Indexes (P1_Col);
```

```
SELECT *
FROM BDF_Table
WHERE S1License_Plate LIKE 'Ford*';
```

The PE will choose to scan the narrow table (the Join Index) and qualify all rows that qualify with a car license LIKE 'Ford*'. Then the PE uses the ROWID to get data from the BDF_Table where row is in the same AMP because both the base table and join index are on the same AMP because they both have the same Primary Index. This can save enormous time for queries using the LIKE command.

A LIKE command on a base table will never use a Non-Unique Secondary Index (NUSI). The above technique should be tested and only used if a lot of users are utilizing the LIKE command on a large table. If that is the case a lot of time can be saved.

Single Table Join Index with Value Ordered NUSI

A Value Ordered NUSI can only be done on columns that are 4-byte integers. Dates qualify because they are stored internally in Teradata as 4-byte integers.

```
CREATE JOIN INDEX Customer AS
SELECT * FROM Order_Table
Primary Indexes (Order_No, Order_Date);

Create Index (Order_Date) order by values
on Order_No;
```

A value ordered index has been expanded from 16 to 64 columns.

Indexes are always sorted by their hash, but a Value Ordered index is sorted on each AMP by values and not hash. This is a great technique for you to run trials on.

Aggregate Join Indexes

Aggregate Join Indexes may be defined on:

Single Tables - A columnar subset of a base table with aggregates automatically maintained by Teradata.


Multiple Tables - A columnar subset of as many as 64 base tables with aggregate columns automatically maintained by Teradata.

Sparse Join Indexes are defined with a WHERE clause that limits the number of base table rows included and the space required to store them.

Aggregate Join Indexes can only include SUM and COUNT values. You can calculate Averages from these two columns though.

Compressed Single-Table Join Index

```
CREATE JOIN INDEX Order_IDX
AS
SELECT (Customer_Number)
      (Order_Number
      (Order_Date
      (Order_Total)
FROM   Order_Table
PRIMARY INDEX (Customer_Number) ;
```



We've duplicated the **Order_Table** with a **different Primary Index**.

This is the compressed version of a Single-Table Join Index. Notice the parentheses around Customer_Number in the SELECT list. A single row is used for each customer, with repeating orders per customer inside the Join Index.

Aggregate Join Index

```
CREATE JOIN INDEX Aggr_Order_IDX AS
SELECT
  Customer_Number
  , Extract(Month from Order_Date) as Yr
  , Extract(Month from Order_Date) as Mo
  , Count(*) as CountYr
  , Sum(Order_Total) as SumYr
FROM   Order_Table
GROUP BY Yr, Mo;
```

Only **Sum** and **Count** can be used with an Aggregate Join Index.

Users never query the Join Index directly. It is the Parsing Engine that commands the AMPs to pull the data from the Join Index. You can extract the Month and Year and calculations are updated as the Base Tables changes. Count and Sum are required to be a data type of FLOAT. Why can you only have Count and Sum? Max and Min aren't that important to know and Average isn't all that important either. Business users want to know how much and how many so AVERAGE, MIN, and MAX are excluded!

New Aggregate Join Index (Teradata V14.10)

```
CREATE JOIN INDEX App_Order_IDX AS
SELECT
  Customer_Number,
  Extract(Year From Order_Date) As Yr,
  Extract(Month From Order_Date) As Mon,
  Country As Country,
  Sum(Order_Total) As Summy,
  Rank (Order_Total) As Rank_Order,
  Min(Order_Total) As Min_Order
FROM   Order_Table
GROUP BY Yr, Mon, Country;
```

Only **Sum** and **Count** can be used with an Aggregate Join Index (pre-Teradata V14.10).

Now, **Min** and **Max** can be used with an Aggregate Join Index (Teradata V14.10).

Aggregate Join Index (AJI) has always supported (SUM and COUNT), but with Teradata V14.10, there is support for MIN/MAX aggregates also.

Sparse Join Index

```
CREATE JOIN INDEX Sparse_Order_IDX AS
SELECT Order_Number,
       Customer_Number,
       Order_Total,
       Order_Date
FROM   Order_Table
WHERE  Order_Date BETWEEN
       '1999-10-01' AND '1999-12-31'
Primary Index (Customer_Number);
```

A **Sparse Join Index** is a Join Index with a **WHERE** Clause!

A Sparse Join Index has a **WHERE** clause so it doesn't take all the rows in the table, but only a portion. This is a very effective way to save space and focus on the latest data.

A Global Multi-Table Join Index

```
CREATE JOIN INDEX EMP_DEPT_IDX AS
SELECT EmpID, EmpName, DeptID, DeptName
FROM   EmpTable, DeptTable
WHERE  EmpID = DeptID;
```

```

-- Department Table as D
-- E:Emps_No, F:Emp_No
-- F:LAST, INDEX (Dept_No)

```

With the ROWID inside the Join Index, the PE can get columns in the User's SQL NOT specified in the Join Index directly from the Base Table by using the Row-ID.

Creating a Hash Index

```

-- CREATE HASH INDEX EMP_HASH_IDX
-- (Dept_No, First_Name, Last_Name)
-- ON EMPLOYEE_TBL;

```

Ordered by Hash of Primary Index

```

-- CREATE HASH INDEX EMP_HASH_VAL
-- (Dept_No, First_Name, Last_Name)
-- ON EMPLOYEE_TBL;
-- ORDER BY VALUES (Dept_No);

```

Ordered by the Value Dept_No

A Hash Index can be Ordered by Values or by Hash.

Join Index Details

- Join Indexes are **physically** stored exactly like normal Teradata tables.
- Users can't query the Join Index directly, but **PE** will decide when to use.
- Join Indexes are **automatically** updated as base tables change.
- Join Indexes can have Non-Unique Primary Indexes (**NUPI**) only.
- Join Indexes can have Non-Unique Secondary (**NUSI**) Indexes.
- Max 64 Columns per Table per Join Index.
- BLOB and CLOB types **cannot** be defined.
- Triggers with Join Indexes allowed V2R6.2.
- After Restoring a Table, Drop and Recreate the Join Index.
- FastLoad**/**MultiLoad** won't load to tables with a Join Index defined.