

CONTENTS:

- 1) The Quantile Function
- 2) Top SQL Command Cheat Sheet
- 3) View Functions

The Quantile Function

The Quantile Function

Overview

"There is fatality in friendship."

William Shakespeare

The Quantile Function and Syntax

The syntax for the Quantile function:

```
QUANTILE (qpartitioner, column_name) ,order-by [DESC | ASC]
[QUALIFY QUANTILE (column_name) (< | > | =) <= | >=] ,num=<of-ranks>}
```

A Quantile is used to divide rows into a number of categories or grouping of roughly the same number of rows in each group. The percentile is the QUANTILE most commonly used in business. The means that the request is based on a value of 100 for the number of partitions. It is also possible to have quantiles (based on 4), tertiles (based on 3) and deciles (based on 10).

By default, both the QUANTILE column and the QUANTILE value itself will be output in ascending sequence. As in some cases, the ORDER BY clause may be used to reorder the output for display. Here, the order of the output does not change the meaning of the output, unlike a summation where the values are being added together and all need to appear in the proper sequence.

A Quantile Example

This example determines the percentile for every row in the Sales table based on the daily sales amount, and sorts it into sequence by the value being categorized, which is Daily_Sales here.

```
SELECT Product_ID, Sales_Date, Daily_Sales
FROM Sales_Table
WHERE Product_ID = 100
ORDER BY Product_ID, Sales_Date, Daily_Sales
```

Product_ID	Sales_Date	Daily_Sales	Quantile
2000	10/02/2000	32000.50	0
2000	10/04/2000	32000.50	0
2000	10/06/2000	34001.99	25
2000	10/07/2000	40001.99	37
2000	10/09/2000	41000.18	50
2000	10/09/2000	54001.10	62
2000	10/01/2000	54000.00	75
2000	10/03/2000	61000.00	87

Notice that the amount of 32000.50 in the first two rows has the same percentile value. They are the same value, and will, therefore, be put into the same partition.

A Quantile Example using DESC Mode

```
SELECT Product_ID, Sales_Date, Daily_Sales
FROM Sales_Table
WHERE Product_ID = 100
ORDER BY Product_ID, Sales_Date, Daily_Sales DESC
```

Product_ID	Sales_Date	Daily_Sales	Quantile
2000	10/02/2000	32000.50	0
2000	10/04/2000	32000.50	10
2000	10/06/2000	34001.99	20
2000	10/07/2000	40001.99	30
2000	10/09/2000	41000.18	40

2000	10/03/2000	43200.18	80
2000	09/29/2000	43950.03	60
2000	10/04/2000	54553.10	70
2000	10/02/2000	54550.25	90
1000	10/03/2000	43700.00	80

Notice the first two rows. This is because the Sale date DESC, impacts the first two rows. Why? Since these rows have the same value, it uses the Sale Date column as a tiebreaker for the sequencing and makes them different from each other. Hence, they are assigned to different values in different partitions.

QUALIFY to find Products in the top Partitions

This example uses a QUALIFY to show only products that sell in the top 60 Percentile

```
SELECT Product_ID, Sale_Date, Daily_Sales
FROM (SELECT Product_ID, Daily_Sales, Sale_Date ) as "Percentile"
QUALIFY "Percentile" >= 60 ;
```

Product_ID	Sale_Date	Daily_Sales	Percentile
2000	09/29/2000	43950.03	60
1000	09/29/2000	43550.40	60
2000	09/29/2000	43550.03	70
1000	09/29/2000	43550.22	70
1000	10/04/2000	54553.10	80
2000	10/02/2000	54550.25	80
2000	09/29/2000	43200.18	90
1000	10/03/2000	43700.00	90

Like the aggregate functions, OLAP functions must read all required rows before performing their operation. Therefore, the WHERE clause cannot be used. Where the aggregates use HAVING, the OLAP functions use QUALIFY. The QUALIFY evaluates the result to determine which ones to return.

QUALIFY to find Products in the top Partitions Sorted DESC

```
SELECT Product_ID, Sale_Date, Daily_Sales
FROM (SELECT Product_ID, Daily_Sales, Sale_Date ) as "Percentile"
ORDER BY "Percentile" DESC ;
```

Product_ID	Sale_Date	Daily_Sales	Percentile
2000	10/03/2000	43700.00	90
2000	09/29/2000	43200.18	90
2000	10/02/2000	54550.25	80
2000	10/04/2000	54553.10	80
2000	09/29/2000	54500.22	70
2000	09/29/2000	43950.03	70

The ORDER BY changes the sequence of the rows being listed, not the meaning of the percentile. The above functions both determined that the highest number in the column is the highest percentile. The data value sequence ascends as the percentile ascends, or descends as the percentile descends. When the sort in the QUANTILE function is changed to ASC, the data value sequence changes to ascend as the percentile descends. The sequence of the percentile does not change. But, the data value sequence is changed to ascend (ASC) instead of the default, which is to descend (DESC).

QUALIFY to find Products in the top Partitions Sorted ASC

```
SELECT Product_ID, Sale_Date, Daily_Sales
FROM (SELECT Product_ID, Daily_Sales ASC, Sale_Date ) as "Percentile"
QUALIFY "Percentile" >= 70 ;
```

Product_ID	Sale_Date	Daily_Sales	Percentile
2000	10/04/2000	32800.50	70
2000	10/07/2000	32800.50	70
3000	10/01/2000	28000.00	80
3000	10/03/2000	21000.75	80
3000	10/02/2000	18479.94	90
3000	10/04/2000	18479.93	90

The example above uses the ASC to cause the data values to go contradictory to the percentile.

QUALIFY to find Products in top Partitions with Tiebreaker

```
SELECT Product_ID, Sale_Date, Daily_Sales
FROM (SELECT Product_ID, Daily_Sales AS D, Sale_Date AS S) AS "Percentile"
QUALIFY "Percentile" >= T0 ;
```

Product_ID	Sale_Date	Daily_Sales	Percentile
2000	10/04/2000	32800.50	70
1000	10/02/2000	32000.50	70
3000	10/01/2000	28000.00	80
3000	10/03/2000	21000.75	80
3000	10/02/2000	18479.94	90
3000	10/04/2000	18479.93	90

The next SELECT modifies the above query to incorporate the sale date as a tiebreaker and reverse the ordering for the top rows with sales of \$32,000.50.

Using Tertiles (Partitions of Four)

```
SELECT Product_ID, Sale_Date, Daily_Sales
FROM (SELECT Product_ID, Daily_Sales AS D, Sale_Date AS S) AS "Quantiles"
WHERE Product_ID in (1000, 2000) ;
```

Product_ID	Sale_Date	Daily_Sales	Quantile
1000	10/04/2000	32000.50	10
2000	10/04/2000	32000.50	10
1000	09/30/2000	30000.07	10
2000	10/02/2000	30000.93	10
1000	10/01/2000	40200.43	10
2000	09/28/2000	41900.00	10
2000	10/03/2000	40200.10	10
1000	09/29/2000	40000.00	10
1000	09/28/2000	40000.00	10
2000	09/30/2000	40000.00	10
1000	09/29/2000	40000.00	10
1000	10/04/2000	34000.10	10
2000	10/02/2000	34000.20	10
1000	10/03/2000	40000.00	10

Instead of 100, the example above uses a quartile (QUANTILE based on 4 partitions).

How Quantile Works

```
SELECT Product_ID, Sale_Date, Daily_Sales
FROM (SELECT Product_ID, Daily_Sales AS D, Sale_Date AS S) AS "Quantiles"
WHERE Product_ID = 1000 ;
```

Product_ID	Sale_Date	Daily_Sales	Quantile
1000	10/02/2000	32000.50	10
1000	09/30/2000	30000.07	10
1000	10/01/2000	40200.43	10

1000	09/28/2000	40150.40	PARTITION
1000	09/28/2000	54500.12	
1000	10/04/2000	54550.10	
1000	10/05/2000	64500.00	

Assigning a different value to the «partitions» indicator of the QUANTILE function changes the number of partitions established. Each Quantile partition is assigned a number starting at 0, increasing to a value that is one less than the partition number specified. So, with a Quantile of 4, the partitions are 0 through 3 and for 10, the partitions are assigned 0 through 9. Then, all the rows are distributed as evenly as possible into each partition from highest to lowest values. Normally, extra rows with the lowest value begin back in the lowest numbered partitions.

Top SQL Commands Cheat Sheet

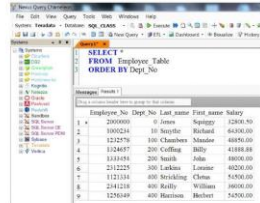
Top SQL Commands Cheat Sheet

Overview

"You miss 100 percent of the shots you never take."

-Vince Lombardi

SELECT All Columns from a Table and Sort



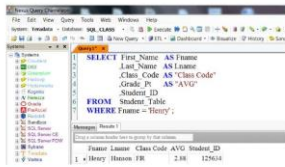
The screenshot shows the Teradata QueryGrid interface. The SQL editor contains the following query:

```
SELECT *  
FROM Employee Table  
ORDER BY Dept_No
```

The Results pane displays the following data:

Employee No	Dept No	Last Name	First Name	Salary
1	20000000	Jones	Scott	12000.00
2	10000000	Smith	John	8000.00
3	12000000	Chambers	William	4000.00
4	11000000	Coffey	Billy	4100.00
5	13000000	Smith	John	8000.00
6	21000000	Levens	Louise	8000.00
7	11000000	Stevens	Chris	5000.00
8	21000000	Reilly	William	8000.00
9	12000000	Harrison	Robert	5000.00

Select Specific Columns and Limiting the Rows




The screenshot shows the Teradata QueryGrid interface. The SQL editor contains the following query:

```
SELECT First Name AS FName,  
Last Name AS LName,  
Class Code AS "Class Code",  
Grade PK AS "AVG"  
FROM Student Table  
WHERE FName = 'Henry';
```


The Results pane displays the following data:

First Name	Last Name	Class Code	AVG
Henry	FR	2.00	1250.00

Changing your Default Database



Keywords that describe you



Select TOP Rows in a Rank Order

SELECT TOP 7 Last_Name,
Class_Code,
Grade_Pt
FROM
Student_Table
ORDER BY Grade_Pt DESC;

Last_Name	Class_Code	Grade_Pt
Thorn	FR	4.00
Boad	SR	3.95
Wilson	SO	2.80
Delaney	SR	3.35
Hillings	SR	3.00
Hanson	FR	2.88
Harris	SO	2.00

A Sample number of rows

SELECT *
FROM Employee_Table
SAMPLE 8

Employee_No	Dept_No	Last_name	First_name	Salary
20000000	0	Jones	Edgney	32400.50
1256400	400	Harrison	Herbert	54500.00
1000214	10	Smythe	Richard	64300.00
1121114	400	Neidking	Clema	54500.00
1212578	100	Chauden	Mauder	48450.00
2341218	400	Rathly	William	36000.00
1121687	200	Coffing	Billy	41888.88
2312325	300	Larkins	Locaine	40200.00

Getting a Sample Percentage of rows

[Find information about a Table](#)

SHOW TABLE Employee_Table

```
CREATE SET TABLE SQL_CLASS.Employee_Table NO FALLBACK,
NO BEFORE JOURNAL,
NO AFTER JOURNAL,
CHECKSUM = DEFAULT
Employee No INHERIT,
Dept_No SMALLINT,
Last_name VARCHAR(25) CHARACTER SET LATIN NOT CASESPECIFIC,
Salary BIGINT(12,2) CHARACTER SET LATIN NOT CASESPECIFIC,
NOLOG PRIMARY INDEX (Employee_No)
INDEX (Dept_No)
```

Using Aggregates

SELECT Dept_No
MIN (Salary) AS "Min"
MAX (Salary) AS "Max"
SUM (Salary) AS "Sum"
AVG (Salary) AS "Avg"
COUNT(*) AS "Cnt"
FROM Employee_Table
GROUP BY Dept_No
ORDER BY 1

Dept_No	Min	Max	Sum	AVG	CNT
1	8	12000.00	12000.00	12000.00	1
2	10	14000.00	14000.00	14000.00	1
3	100	18000.00	18000.00	18000.00	1
4	200	21000.00	21000.00	21000.00	2
5	300	20000.00	20000.00	20000.00	1
6	400	20000.00	20000.00	20000.00	2

Performing a Join

Performing a Join using ANSI Syntax

Using Date, Time and Timestamp

The screenshot shows the Oracle SQL Developer application. The top menu bar includes File, Edit, View, Query, Tools, View, Windows, and Help. The top toolbar contains icons for various actions like opening files, saving, and executing queries. The main window is titled 'New Query Connection' and shows a tree view on the left with 'Database' selected. The central pane displays a SQL query:

```
SELECT Date
      ,Current Date
      ,Time
      ,Current Time
      ,Current Timestamp(5)
```

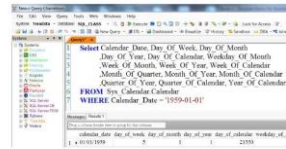
Below the query, the 'Messages' pane shows 'Results: 1' and a message: 'Drop column before here to group by this column.' The 'Results' pane displays a table with the following data:

Date	Time	Current Date	Current Time	Current Timestamp(5)
09/13/2013	15:10:04	09/13/2013	15:10:04	09/13/2013 15:10:04.00000

Using Date Functions

[illegible]

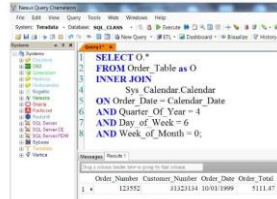
Using the System Calendar



```
SELECT Calendar.Date, Day_Of_Week, Day_Of_Month,
       Day_Of_Year, Day_Of_Calendar, Weekday_Of_Month,
       Week_Of_Month, Week_Of_Year, Week_Of_Calendar,
       Month_Of_Quarter, Month_Of_Year, Month_Of_Calendar,
       Quarter_Of_Year, Quarter_Of_Calendar, Year_Of_Calendar
FROM Sys_Calendar.Calendar
WHERE Calendar.Date = '1979-01-01'
```

Results: Row(s): 1
Col(s): 12
Calendar.Date, Day_Of_Week, Day_Of_Month, Day_Of_Year, Day_Of_Calendar, Weekday_Of_Month, Week_Of_Month, Week_Of_Year, Week_Of_Calendar, Month_Of_Quarter, Month_Of_Year, Month_Of_Calendar
1 * 01/01/1979 0 1 1 22379

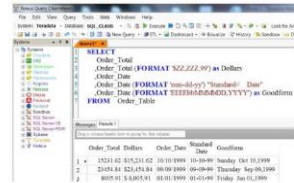
Using the System Calendar In a Query



```
SELECT O *
FROM Order_Table as O
INNER JOIN
      Sys_Calendar.Calendar
ON Order_Date = Calendar.Date
AND Quarter_Of_Year = 4
AND Day_of_Week = 6
AND Week_of_Month = 0;
```

Results: Row(s): 1
Col(s): 13
Order_Number, Customer_Number, Order_Date, Order_Total
1 * 12352 31125114 10/01/1999 5111.47

Formatting Data



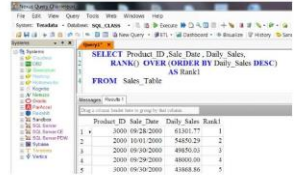
The screenshot shows the Teradata QueryGrid interface. The SQL editor contains the following query:

```
SELECT  
  Order_Table  
  Order_Totals (FORMAT 'ZZZZZ.99') as Dollars  
  Order_Date  
  Order_Date (FORMAT 'mm-dd-yy') as Standard Date  
  Order_Date (FORMAT 'YYYYMMDD.VVYY') as Goodform  
FROM Order_Table
```

The Results pane shows the following data:

Order_Totals Dollars	Order_Date	Standard Date	Goodform
1	15231.62	5/15/2011	20110515.11
2	21451.81	5/21/2011	20110521.11
3	8801.01	5/16/2011	20110516.11

Using Rank



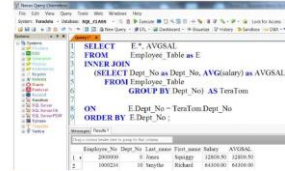
The screenshot shows the Teradata QueryGrid interface. The SQL editor contains the following query:

```
SELECT Product_ID, Sales, Date, Daily_Sales,  
  RANK() OVER (ORDER BY Daily_Sales DESC)  
  AS Rank1  
FROM Sales_Table
```

The Results pane shows the following data:

Product_ID	Sales	Date	Daily_Sales	Rank1
1	3000	05-28-2000	61301.77	1
2	2000	10-01-2000	54650.23	2
3	2000	09-30-2000	48650.03	3
4	2000	09-29-2000	48000.00	4
5	2000	09-28-2000	43888.86	5

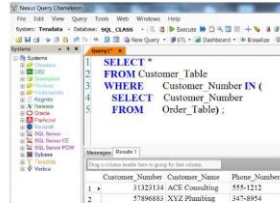
Using a Derived Table



```
1 SELECT E*, AVG(SAL)
2 FROM
3   EMPLOYEES Table as E
4 INTERSECT
5   (SELECT Dept_No as Dept_No, AVG(salary) as AVG(SAL)
6    FROM EMPLOYEES Table
7    GROUP BY Dept_No) AS TeraTom
8 ON
9   E.Dept_No = TeraTom.Dept_No
10 ORDER BY E.Dept_No;
```

Employee_No	Dept_No	Last_name	First_name	Salary	AVG(SAL)
1	200000	Jones	Walter	12000.00	12000.00
2	100000	Neely	Richard	84000.00	84000.00

Using a Subquery



```
1 SELECT *
2 FROM Customer Table
3 WHERE Customer Number IN (
4   SELECT Customer Number
5   FROM Order_Table);
```

Customer Number	Customer Name	Phone Number
1	31325134 ACE Consulting	554-1212
2	57866883 XYZ Plumbing	347-8954

Correlated Subquery

The screenshot shows a Teradata QueryGrid window with a SQL query and its results. The query is as follows:

```

SELECT
FROM Employee_Table as EE
WHERE Salary < 1
SELECT AVG(Salary)
FROM Employee_Table as EEE
WHERE EE.Dept_No = EEE.Dept_No);

```

The results are displayed in a table with the following columns: Employee_No, Dept_No, Last_Name, First_Name, Salary.

Employee_No	Dept_No	Last_Name	First_Name	Salary
1	1223456	Smith	John	48000.00
2	1226440	Harrison	Robert	14500.00
3	1223344	King	James	14500.00

Using Substring

The screenshot shows a Teradata QueryGrid window with a SQL query using the SUBSTRING function. The query is as follows:

```

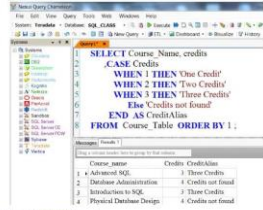
SELECT First_Name,
SUBSTRING(First_Name FROM 2 for 3) AS Quiz
FROM Employee_Table;

```

The results are displayed in a table with the following columns: First_Name, Quiz.

First_Name	Quiz
John	ohn
Robert	ober
James	ames
Michael	icha
Stefan	tefan
Shelley	helly

Basic CASE Statement

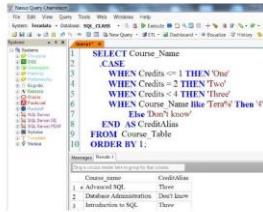


```
SELECT Course_Name, credits
FROM Course_Table
ORDER BY 1;
```

Results (Results 1)

Course_name	Credits	CreditsAlias
1. Advanced SQL	3	Three Credits
2. Database Administration	4	4 Credits not found
3. Introduction to SQL	3	Three Credits
4. Physical Database Design	4	4 Credits not found

Advanced CASE Statement



```
SELECT Course_Name
FROM Course_Table
ORDER BY 1;
```

Results (Results 1)

Course_name	CreditsAlias
1. Advanced SQL	Three
2. Database Administration	Three
3. Introduction to SQL	Three

Using an Access Lock in your SQL

The screenshot shows the Teradata Query Window with the following SQL query:

```
LOCKING ROW FOR ACCESS
SELECT * FROM Employee_Table
```

The Results pane displays the following data:

Employee_No	Dept_No	Last_name	First_name	Salary
1	2000000	Jones	Spartan	22800.00
2	1100000	Smith	John	48000.00
3	1000214	Smith	Richard	64000.00
4	1256149	Harrison	Herbert	54500.00
5	1222778	Chambers	Mandie	48500.00
6	1121334	Stevenson	Christa	54500.00

Collect Statistics

The screenshot shows the Teradata Query Window with the following SQL command:

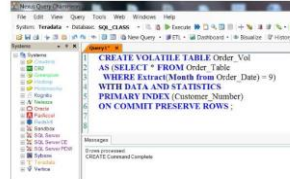
```
COLLECT STATISTICS
COLUMN (Employee_No)
ON Employee_Table;

COLLECT STATISTICS
ON Employee_Table
COLUMN Dept_No;
```

The Messages pane shows the following output:

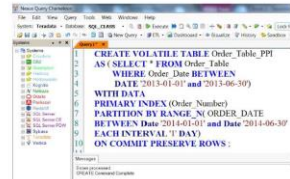
```
Table processed
Column processed
Table processed
Column processed
```

CREATING a Volatile Table with a Primary Index



```
1 CREATE VOLATILE TABLE Order_Vol
2 AS (SELECT * FROM Order_Table
3 WHERE Extract(Month from Order_Date) = 9)
4 WITH DATA AND STATISTICS
5 PRIMARY INDEX (Customer_Number)
6 ON COMMIT PRESERVE ROWS;
```

CREATING a Volatile Table that is Partitioned (PPI)



```
1 CREATE VOLATILE TABLE Order_Table_PPI
2 AS (SELECT * FROM Order_Table
3 WHERE Order_Date BETWEEN
4 DATE '2013-01-01' and '2013-06-30')
5 WITH DATA
6 PRIMARY INDEX (Order_Number)
7 PARTITION BY RANGE_N (ORDER_DATE
8 BETWEEN Date '2014-01-01' and Date '2014-06-30'
9 EACH INTERVAL 1 DAY)
10 ON COMMIT PRESERVE ROWS;
```

CREATING a Volatile Table that is deleted after the Query

Finding out how much Space you have

Query Window: SQL Editor

```

SELECT UserNm,
       CreateName,
       PermSpace,
       SpoolSpace,
       TempSpace,
       LastAlterTimeSp
FROM DBC.Users
WHERE UserNm = USER;

```

Results:

UserNm	CreateName	PermSpace	SpoolSpace	TempSpace	LastAlter
DBRC	DBRC	271189582.56	385865162.56	385865162.56	02/08/2011

How much Space you have Per AMP

Query Window: SQL Editor

```

LOCKING ROW FOR ACCESS
SELECT Types AS "AMPs",
       DatabaseName AS "Database",
       AccountName AS "Account",
       (MaxPerm) AS "Perm",
       (MaxSpool) AS "Spool",
       (MaxTemp) AS "Temp"
FROM DBC.DatSpace
WHERE DatabaseName = 'sql01';

```

Results:

AMPs	Database	Account	Perm	Spool	Temp
1	SQL01	DBRC	2500000.00	10000000.00	500000.00
2	SQL01	DBRC	2500000.00	10000000.00	500000.00

Finding your Space

```

1 LOCKING ROW FOR ACCESS
2 SELECT SUM(MaxPerm) AS "Perm"
3       ,SUM(MaxSpool) AS "Spool"
4       ,SUM(MaxTemp) AS "Temp"
5 FROM DBC.DiskSpace
6 WHERE DatabaseName = USER;

```

Messages: Results 1

Perm	Spool	Temp
271140562.56	385563682.56	385563682.56

Finding Space Skew in Tables in a Database

```

1 SELECT
2   Vproc
3   ,CAST (TableName AS CHAR(20))
4   ,CurrentPerm
5   ,PeakPerm
6 FROM DBC.TableSizeV
7 WHERE DatabaseName = 'SQL_Class'
8 ORDER BY TableName, Vproc ;

```

Messages: Results 1

TableName	Vproc	CurrentPerm	PeakPerm
Address	1	1768.00	1776.00
Address	2	1824.00	1824.00
AddressTLC	3	17920.00	17920.00
AddressTLC	4	17920.00	17920.00

Finding the Number of rows per AMP for a Column

The screenshot shows the Teradata Query Window with the following SQL query:

```
LOCKING ROW FOR ACCESS SELECT  
HASHBAMP (HASHBUCKET  
(HASHROW (Employee_No))) AS "AMP #"  
,COUNT(*)  
FROM SQL_Class.Employee_Table  
GROUP BY 1  
ORDER BY 1
```

The results pane shows the output of the query:

AMP #	Count(*)
1	3
2	6

Finding Account Information

The screenshot shows the Teradata Query Window with the following SQL query:

```
SELECT *  
FROM USC.AccountInfoV  
ORDER BY 1;
```

The results pane shows the output of the query:

UserName	AccountName	UserProfile
1 + AdminProfile	ADMIN12	Profile
2 + CaringCare	UNC	User
3 + Compagnard	DMComp	User
4 + comale	Bitronics-comale-us	User
5 + Creditbump	Creditbump	User

Ordered Analytics



capgemini

View Functions

View Functions

Overview

"It is easier to go down a hill than up it, but the view is much better at the top."

- Arnold Bennett

Creating a Simple View

```
Employee_Table
Employee_No  Dept_No  Last_Name  First_Name  Salary
-----
10000000    10       Jones      'Eugene'    12000.00
10001234    10       Smythe     Richard    32000.00
10001235    200       Chambers   Andrew     40000.00
10001236    200       Coffey     Billy      41000.00
10001237    200       Smith      John       40000.00
10001238    200       Jackson    William    30000.00
10001239    400       Harrison   Robert     54000.00
10001240    400       Bailey     William    50000.00
10001241    400       Harrington Clinton  54000.00
```

```
CREATE View Employee_V AS
SELECT  EMPLOYEE_NO,
        DEPT_NO,
        FIRST_NAME,
        LAST_NAME,
        SALARY
FROM    Employee_Table_1
```

The purposes of views are to restrict access to certain columns, derive columns or Join Tables, and to restrict access to certain rows (if a WHERE clause is used).

Basic Rules for Views

1. No **ORDER BY** inside the View CREATE (some exceptions exist)
2. All Aggregation needs to have an **ALIAS**
3. Any Derived columns (such as Math) needs an **ALIAS**


```
CREATE View Department_Salaries AS
SELECT  Dept_No,
        SUM(Salary) as SumSal
FROM    Employee_Table
GROUP BY Dept_No
```

```
SQL >
FROM Department_Salaries
ORDER BY 1,2
```

```
Dept_No  SumSal  MonthSal
-----
10       12000.00  2700.00
20       161000.00  5366.66
400      148000.00  4933.33
```

Above, are the basic rules of Views with excellent examples.

How to Modify a View



The REPLACE Keyword will allow a user to change a view.

Exceptions to the ORDER BY Rule inside a View



There are EXCEPTIONS to the ORDER BY rule. The TOP command allows a view to work with an ORDER BY inside. ANSI OLAP statements also work inside a View.

How to Get HELP with a View

1

HELP View Command

HELP View Emp_View ;

Column Name	Table	Comment	Nullable	Format	Title
Employee_No	?		?	?	?
Dept_No	?		?	?	?
Last_Name	?		?	?	?
First_Name	?		?	?	?
Salary	?		?	?	?

Row Length	Column Total Width	Positional Position Data	Rowset Size	Rowset Width
?	?	?	?	?
?	?	?	?	?
?	?	?	?	?
?	?	?	?	?

Input Type	Table View?	Default Collate	Char Type	IM Coll Type
?	?	?	?	?
?	?	?	?	?
?	?	?	?	?

The Help View command does little but show you the columns.

Views sometimes CREATED for Formatting or Row Security

```
CREATE VIEW emp1_200_v AS
SELECT Employee_No AS Emp_No
       Last_Name AS Last_Name
       Salary/12 (Column '000,000.00') AS Monthly_Salary
FROM Employee_Table
WHERE Dept_No = 200;
```

SELECTING from a View

```
SELECT *
FROM emp1_200_v
ORDER BY Monthly_Salary;
```

Emp_No Last_Name Monthly_Salary

123456 Smith 24,000.00

Views are designed to do many things. In the example above, this view formats and derives data, limits columns, and also limits the rows coming back with a WHERE.

Another Way to Alias Columns in a View CREATE

```
CREATE VIEW emp1_200_v (Emp_No, Last_Name, Monthly_Salary)
AS
SELECT Employee_No
       Last_Name
       Salary/12 (Format '$555,555.99')
FROM Employee_Table
WHERE Dept_No = 200;
```

Shows me here

```
SELECT *
FROM Emp1_200_v
ORDER BY Monthly_Salary;
```

Emp_No	Last_Name	Monthly_Salary
1121457	Coelho	24,400.78
1121454	Smith	24,000.00

Will this View CREATE Error? Not! It won't error because it's Aliased above!

Resolving Aliasing Problems in a View CREATE

```
CREATE VIEW emp1_200_v (Emp_No, Last, Monthly_Salary)
AS SELECT Employee_No
      Last_Name
      Salary/12 (format '$$$$,$$9.99') as Sal_Monthly
FROM Employee_Table
WHERE Dept_No = 200;
```

```
SELECT *
FROM Emp1_200_v
ORDER BY Sal_Monthly;
```

Emp_No	Last_Name	Monthly_Salary	First Alias
1121457	Coelho	\$2,400.78	
1121454	Smith	\$4,000.00	

The ALIAS for Salary/12 that'll be used in this example is MONTHLY_SALARY. It came first at the top, even though it is aliased in the SELECT list also.

Resolving Aliasing Problems in a View CREATE

```
CREATE VIEW emp1_200_v (Emp_No, Last, Monthly_Salary)
AS SELECT Employee_No
      Last_Name
      Salary/12 (format '$$$$,$$9.99') as Sal_Monthly
FROM Employee_Table
WHERE Dept_No = 200;
```

```
SELECT *
FROM Emp1_200_v
ORDER BY Sal_Monthly;
```

What will happen in the above query?

Resolving Aliasing Problems in a View CREATE

```
CREATE VIEW emp1_200_v (Emp_No, Last, Monthly_Salary)
AS SELECT Employee_No
      Last_Name
      Salary/12 (format '$$$$,$$9.99') as Sal_Monthly
FROM Employee_Table
```

```
WHERE Dept_No = 200 ;
```

```
SELECT *
FROM Empl_200_v
ORDER BY Sal_Xtenfly ;
```

Doesn't Recognize

ERROR

If you ALIAS at the top, then that is the only ALIAS that the query can recognize. So, it is a good idea to alias at the top or the bottom, but not do both.

CREATING Views for Complex SQL such as Joins

```
CREATE VIEW Customer_Order_v AS
SELECT Customer_Name AS Customer,
       Order_Number AS Order_Number,
       (Order_Total / 1000000) AS Total_Amount
FROM Customer_Table AS Cust
       Order_Table AS Ord
WHERE Cust.Customer_Number = Ord.Customer_Number ;

SELECT * FROM Customer_Order_v
ORDER BY 2 ;
```

Customer	Order_Number	Total_Amount
AW Consulting	123456	78.11149
Billy's Best Choice	123456	612.344.93
Billy's Best Choice	123456	94.009.91
Dadabene B-C	123456	215.231.42
EDC Consulting	123777	421.474.14

A huge reason for Views, other than security, is to make Complex SQL easy for users. This view already has the inner join built into it, but users just SELECT.

WHY certain columns need Aliasing in a View

```
CREATE VIEW Aggreg_Order_v AS
SELECT Customer_Number
       ,Order_Total / 1000000 (format '9999-99') AS Tr_Mth_Order
       ,COUNT(Order_Total) AS Order_Cnt
       ,SUM(Order_Total) AS Order_Sum
       ,AVG(Order_Total) AS Order_Avg
FROM Order_Table ;
```

Customer_Number	Order_Sum
1122114	1111.49
123456	1231.42
1111111	8005.91
1111111	1247.93
123456	2145.14

When you CREATE a view, you have to ALIAS any aggregation or derived data (such as math). Why? So you can SELECT it later without having to do a SELECT *. Here, we only chose two columns and used their ALIAS to retrieve them.

Aggregates on View Aggregates

```
CREATE VIEW Aggregates_Order_v AS
SELECT Customer_Number,
       Order_No/100+100000 (format '9999-99') AS Tr_Mth_Order,
       count(Order_No) AS Order_Cnt,
       sum(Order_Total) AS Order_Sum,
       avg(Order_Total) AS Order_Avg
FROM Order_Table
GROUP BY Customer_Number, Tr_Mth_Order;
```

```
SELECT Customer_Number,
       Order_No
FROM Aggregates_Order_v;
```

Customer_Number	Order_No
11111111	1111-1
11111111	1111-2
11111111	1111-3
11111111	1111-4
11111111	1111-5

```
SELECT sum (Order_Sum)
FROM Aggregates_Order_v;
```

sum (Order_Sum)
64131.3

The examples above show how we put a SUM on the aggregate Order_Sum.

Locking Row for Access

Lock
Placed

```
CREATE VIEW Emp_HR_v AS
LOCKING ROW FOR ACCESS
SELECT Employee_No,
       Dept_No,
       Last_Name,
       First_Name
FROM Employee_Table;
```

```
SELECT * FROM Emp_HR_v;
```

The Employee_Table used above will automatically use an ACCESS Lock, which allows ACCESS during UPDATES or table loads.

Most views utilize the Locking row for ACCESS command. This is because they want to be able to read while a table is being updated and loaded into. If the user knows a dirty read won't have a huge effect on their job, why not make a view lock with an ACCESS Lock, thus preventing unnecessary waiting?

Creating Views for Temporal Tables

```
CREATE VIEW sq001.Props_Ar_2a
AS
LOCKING ROW FOR ACCESS
CREATE TEMPORARY TABLE
SELECT Date_B,
       Prop_No,
       Bldg (Prop_Val_Time) AS Bldg_Val_Time;
```

```
CREATE VIEW sq001.Props_Ar_Max
AS
LOCKING ROW FOR ACCESS
NOINDEPENDENT VALIDATION
SELECT Date_B,
       Prop_No,
       Bldg (Prop_Val_Time) AS Bldg_Val_Time;
```

```
END(FEmp_Vol_Time) AS End_Vol_Time,  
FROM FEmpVol_Connec);  
SELECT * FROM SQL01.FEmp_As_Ts_1;  
SELECT * FROM SQL01.FEmp_As_Max_1;
```

You can create views that will allow users to see the way things are or the way things were. Above, are two excellent examples.

Altering a Table

```
CREATE VIEW Emp_AB_v AS  
SELECT Emp_Vol_ID,  
       Dept_ID,  
       Last_Name,  
       First_Name  
FROM Employee_Table_1;
```

```
-- Altering the actual Table  
ALTER TABLE Employee_Table  
ADD Mgr INT(12);
```

↓

Will the View

STILL run?

```
-- Select from the View  
SELECT * FROM Emp_AB_v;
```

↓

This view will run after the table has added an additional column!

Altering a Table after a View has been created

```
CREATE VIEW Emp_AB_v4 AS  
SELECT *  
FROM Employee_Table1;
```

```
-- Altering the actual Table  
ALTER TABLE Employee_Table4  
ADD Map_No INTEGER;
```

Will the View STILL run?

```
-- Select from the View  
SELECT * FROM Emp_HR_v4;
```

YES!

This view runs after the table has added an additional column, but it won't include Map_No in the view results even though there is a SELECT * in the view. The View includes only the columns present when the view was CREATED.

A View that errors After an ALTER

```
CREATE VIEW Emp_HR_v5 AS  
SELECT Emp_ID,  
       Emp_Name,  
       Emp_Surname  
FROM Employee_Table4;
```

```
-- Altering the actual Table  
ALTER TABLE Employee_Table4  
DROP Emp_ID;
```

Will the View STILL run?

```
-- Select from the View  
SELECT * FROM Emp_HR_v5;
```

ERROR

This view will NOT run after the table has dropped a column referenced in the view.

Troubleshooting a View

```
CREATE VIEW Emp_HR_v6 AS  
SELECT *  
FROM Employee_Table4;
```



This view will NOT run after the table has dropped a column referenced in the view, even though the View was CREATED with a SELECT *. At View CREATE Time, the columns present were the only ones the view considered responsible for. Dept_No was one of those columns. Once Dept_No was dropped, the view no longer works.

Updating Data in a Table through a View

```
CREATE VIEW Emp_HR_v6 AS
SELECT *
FROM Employee_Table6;
```

```
--Updating the table through the View
UPDATE Emp_HR_v6
SET Salary = 8888.88
WHERE Employee_No = 2000000;
```

```
--SELECT from the actual Table
SELECT *
FROM Employee_Table6
WHERE Employee_No = 2000000;
```

Employee_No	Dept_No	Last_Name	First_Name	Salary
2000000	1	Stevens	Equipez	8888.88

You can UPDATE a table through a View if you have the RIGHTS to do so.

Maintenance Restrictions on a Table through a View

There are a few restrictions that disallow maintenance activity on a view with an INSERT, UPDATE or DELETE request. A view cannot be used for maintenance if it:

1. Performs a join operation - more than one table
2. Selects the same column twice - wouldn't know which one to use
3. Derives data - because it does not undo the math or calculation
4. Performs aggregation - because this eliminates detail data
5. Uses OLAP functions - because OLAP data is calculated
6. Uses a DISTINCT or GROUP BY - eliminates duplicate rows

Perform maintenance on a table through a view, but see the restrictions above first.